

DVS Technologies

Compiled and Scrutinized by
Mr. Shaan Shaik
(Senior DevOps Lead)

Words To The Students

Though we have taken utmost efforts to present you this book error free, but still it may contain some errors or mistakes. Students are encouraged to bring, if there are any mistakes or errors in this document to our notice. So that it may be rectified in the next edition of this document.

“Suppressing your doubts is Hindering your growth”.

We urge you to work hard and make use of the facilities we are providing to you, because there is no substitute for hard work. We wish you all the best for your future.

“The grass isn’t greener on the other side; the grass is greener where you water it.”

You and your suggestions are valuable to us; Help us to serve you better. In case of any suggestions, grievance, or complaints, please feel free to write us your suggestions, grievance and feedback on the following

Dvs.training@gmail.com

1. Introduction to Kubernetes

What is Kubernetes:

Commonly referred to by its internal designation during the development as Google(k8s),Kubernetes is an open source container cluster manager.

When it was released in July 2015, Google donated it as "seed technology" to the newly formed Cloud native computing foundation established in a partnership with the Linux Foundation.

Kubernetes primary goal is to provide a platform for automating deployment, scaling and operations of application containers across a cluster of hosts.

Design Overview:

Kubernetes is built through the definition of a set of components (building blocks or "primitives") which, when used collectively, provide a method for the deployment, maintenance and scalability of container based application clusters.

These "primitives" are designed to be loosely coupled (i.e. where little to no knowledge of the other component definitions is needed to use) as well as easily extensible through an API. Both the internal components of Kubernetes as well as the extensions and containers make use of this API.

Although Kubernetes was designed and used internally at Google to deploy and utilize "billions of containers", since it was open sourced under the Apache common license, it has since been adopted formally as a service available from each of the major cloud providers.

Components:

Building Block of Kubernetes:

- Node (master & workers)**
- Pods**
- Labels**
- Selectors**
- Controllers/Deployments**
- Services**
- Control Plane**
- API**

Understanding Kubernetes Architecture :

Explain in detail about the architecture of K8s how it works.

Please check "HighLevel Architecture" figure for more information ...

From the above architecture we can conclude that we will have a "master controller" which will control all the other nodes. Here in k8s we called these nodes as "worker nodes", which means that each and every server (worker nodes) should have docker package installed. Along with this we are going to create our containers in separate environment which we called it as "Pods". Hence you are not allowed to create as separate entity. You should create the containers as part of the pods. If you want to increase the containers you can increase the containers as part of that pods. And finally all the worker nodes will be managed by our master controller.

Components of Kubernetes in detail over view :

worker nodes (Nodes):

You can think of these as "container clients". These are the individual hosts(physical or virtual) that docker is installed on and hosts the various containers within your managed cluster. So Finally "worker nodes" are nothing but our physical or virtual server where we have our container package installed and it acts like a slave to the master controller.

Pods:

A pod consists of one or more containers. Those containers are guaranteed (by the cluster controller) to be located on the same host machine in order to facilitate sharing of resources. Pods are assigned unique ips within each cluster. These allow an application to use ports without having to worry about conflicting port utilization.

Pods can contain definitions of disk volumes or share, and then provide access from those to all the members(containers) within the pod.

Finally, pod management is done through the API or delegated to a controller.

Labels:

Clients can attach "key-value pairs" to any object in the system (like pods or worker nodes).

These become the labels that identify them in the configuration and management of them.

Selectors:

Label selectors represent queries that are made against those labels. They resolve to the corresponding matching objects.

These two item are the primary way that grouping is done in kubernetes and determine which components that a given operation applies to when indicated.

Controllers:

These are used in the management of your cluster. Controllers are the mechanism by which

your desired configuration state is enforced. Controllers manage a set of pods and depending on the desire configuration state, may engage other controllers to handle replication and scaling (Replication controller)

of XX number of containers and pods across the cluster. It is also responsible for replacing any container in a pod that fails (based on the desired state of the cluster). Other controllers that can be engaged include a DaemonSet Controller (enforces a 1 to 1 ratio of pods to worker nodes) and job Controller(that runs pods to "completion", such as in batch jobs).

Each set of pods any controller manages, is determined by the label selectors that are part of its definition.

Services:

A pod consists of one or more containers. Those containers are guaranteed (by the cluster controller) to be located on the same host machine in order to facilitate sharing of resources. This is so that pods can work together, like in a multi-tier application configuration. Each set of pods that define and implement a service (like Mysql or Apache) are defined by the label selector.

Kubernetes can then provide service discovery and handle routing with the static ip for each pod as well as load balancing (round robin based) connections to that service among the pods that match the label selector indicated. By default although a service is only exposed inside a cluster, it can also be exposed outside a cluster as needed

On the Master Node following components will be installed :

API Server – It provides kubernetes API using Json / Yaml over http, states of API objects are stored in etcd

Scheduler – It is a program on master node which performs the scheduling tasks like launching containers in worker nodes based on resource availability

Controller Manager – Main Job of Controller manager is to monitor replication controllers and create pods to maintain desired state.

etcd – It is a Key value pair data base. It stores configuration data of cluster and cluster state.

Kubectl utility – It is a command line utility which connects to API Server on port 6443. It is used by administrators to create pods, services etc.

On Worker Nodes following components will be installed ::

Kubelet – It is an agent which runs on every worker node, it connects to docker and takes care of creating, starting, deleting containers.

Kube-Proxy – It routes the traffic to appropriate containers based on ip address and port number of the incoming request. In other words we can say it is used for port translation.

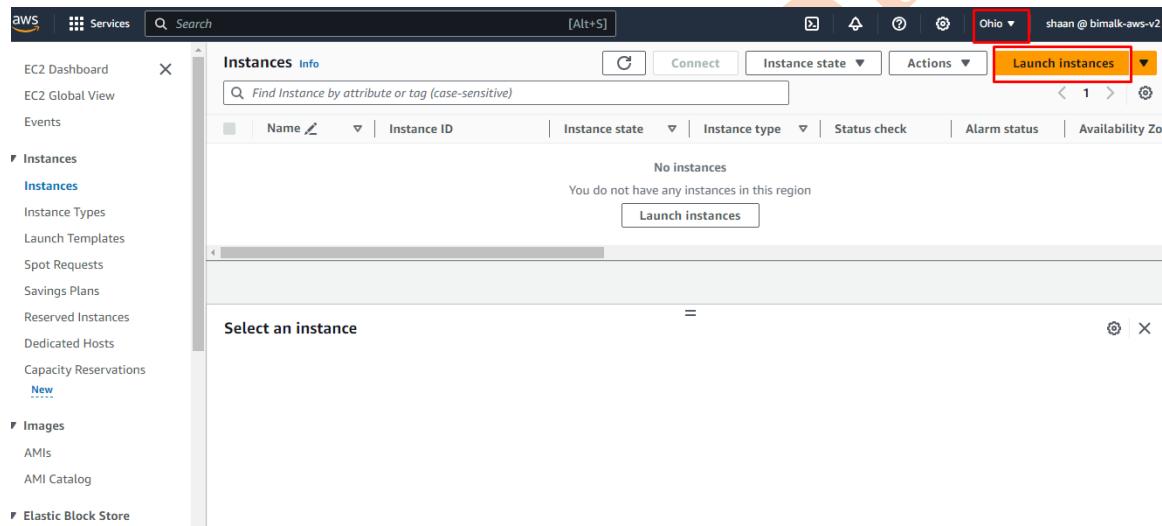
Pod – Pod can be defined as a multi-tier or group of containers that are deployed on a single worker node or docker host.

2 Preparing workstation

Read the below before you start (mandatory):

- 1 Since it's a workstation we can go ahead and create workstation with t2.micro
- 2 Make sure that you are configuring worker node groups with t2.medium configuration
- 3 Note that all the servers are getting created in the same network(vpc)
for example in future jenkins (t2.medium) etc ..
- 4 Make sure you are creating the infrastructure in OHIO (us-east-2)

1. Creating workstation:



Name and tags [Info](#)

Name [Add additional tags](#)

Application and OS Images (Amazon Machine Image) [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below.

Quick Start

Amazon Linux aws macOS Ubuntu Windows Red Hat SUSE L

Mac **ubuntu®** **Microsoft** **Red Hat** **SUSE**

Browse more AMIs Including AMIs from AWS, Marketplace and the Community

Amazon Machine Image (AMI)

Amazon Linux 2 AMI (HVM) - Kernel 5.10 SSD Volume Type ami-0f599bbc07afc299a (64-bit (x86)) / ami-04d758931d5479131 (64-bit (Arm)) Virtualization: hvm ENA enabled: true Root device type: ebs	Free tier eligible
--	--------------------

Summary

Number of instances [Info](#)
1

Software Image (AMI)
Amazon Linux 2023 AMI 2023.2.2... [read more](#)
ami-06dd4b7182ac3480fa

Virtual server type (instance type)
t2.micro

Firewall (security group)
New security group

Storage (volumes)
1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 30 GiB of EBS storage, 2 million I/Os, 1 GB of snapshots, and 100 GB of bandwidth to the Internet.

[Launch instance](#) [Review commands](#)

Instance type

t2.micro
Family: t2 1 vCPU 1 GiB Memory Current generation: true
On-Demand Linux base pricing: 0.0116 USD per Hour
On-Demand SUSE base pricing: 0.0116 USD per Hour
On-Demand Windows base pricing: 0.0162 USD per Hour
On-Demand RHEL base pricing: 0.0716 USD per Hour

A additional cost applies for AMIs with pre-installed software.

Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name required [Create new key pair](#)

Network settings [Info](#)

Network [Info](#)
vpc-0f651a58721284131

Subnet [Info](#)
No preference (Default subnet in any availability zone)

Auto-assign public IP [Info](#)
Enable

Summary

Number of instances [Info](#)
1

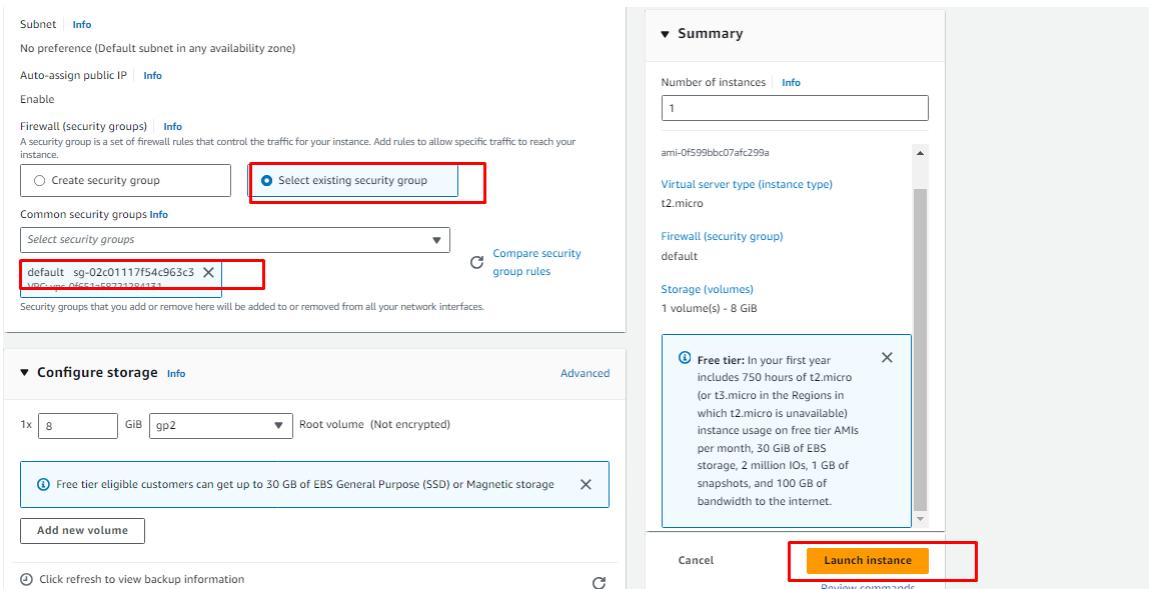
Virtual server type (instance type)
t2.micro

Firewall (security group)
New security group

Storage (volumes)
1 volume(s) - 8 GiB

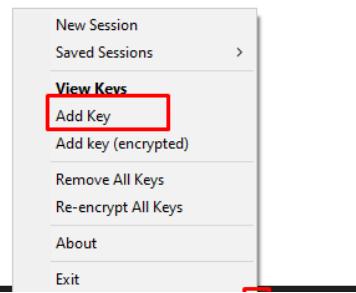
Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 30 GiB of EBS storage, 2 million I/Os, 1 GB of snapshots, and 100 GB of bandwidth to the Internet.

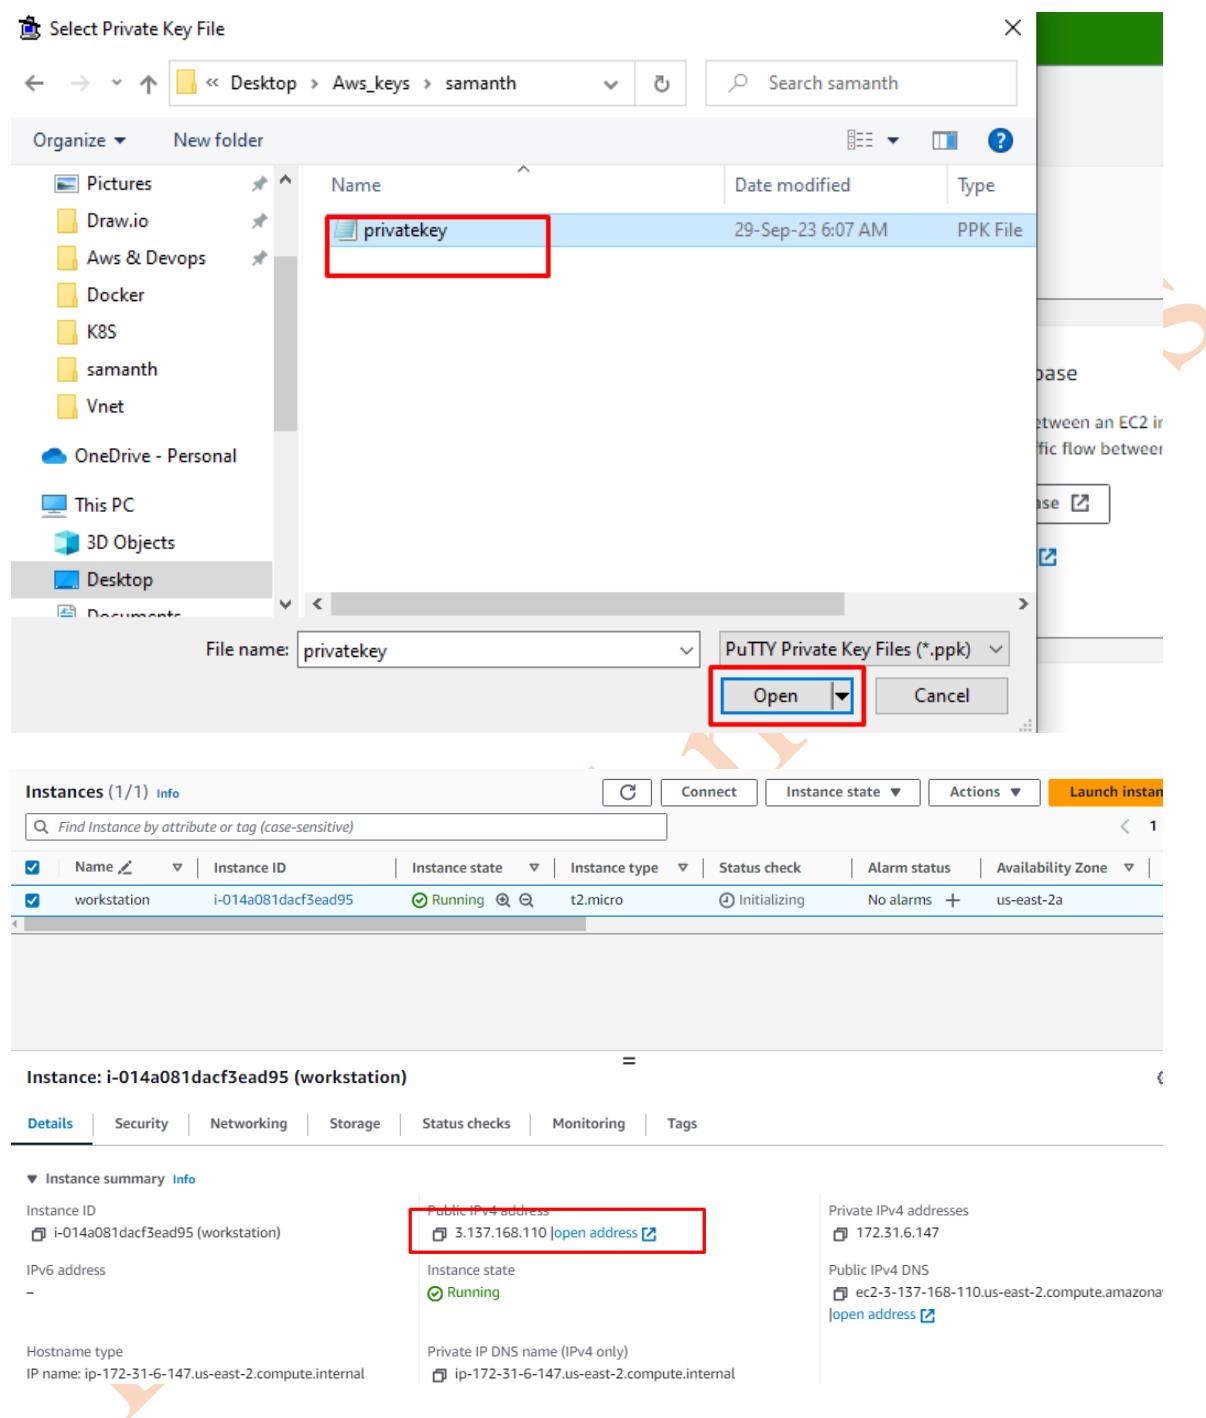
[Launch instance](#) [Review commands](#)

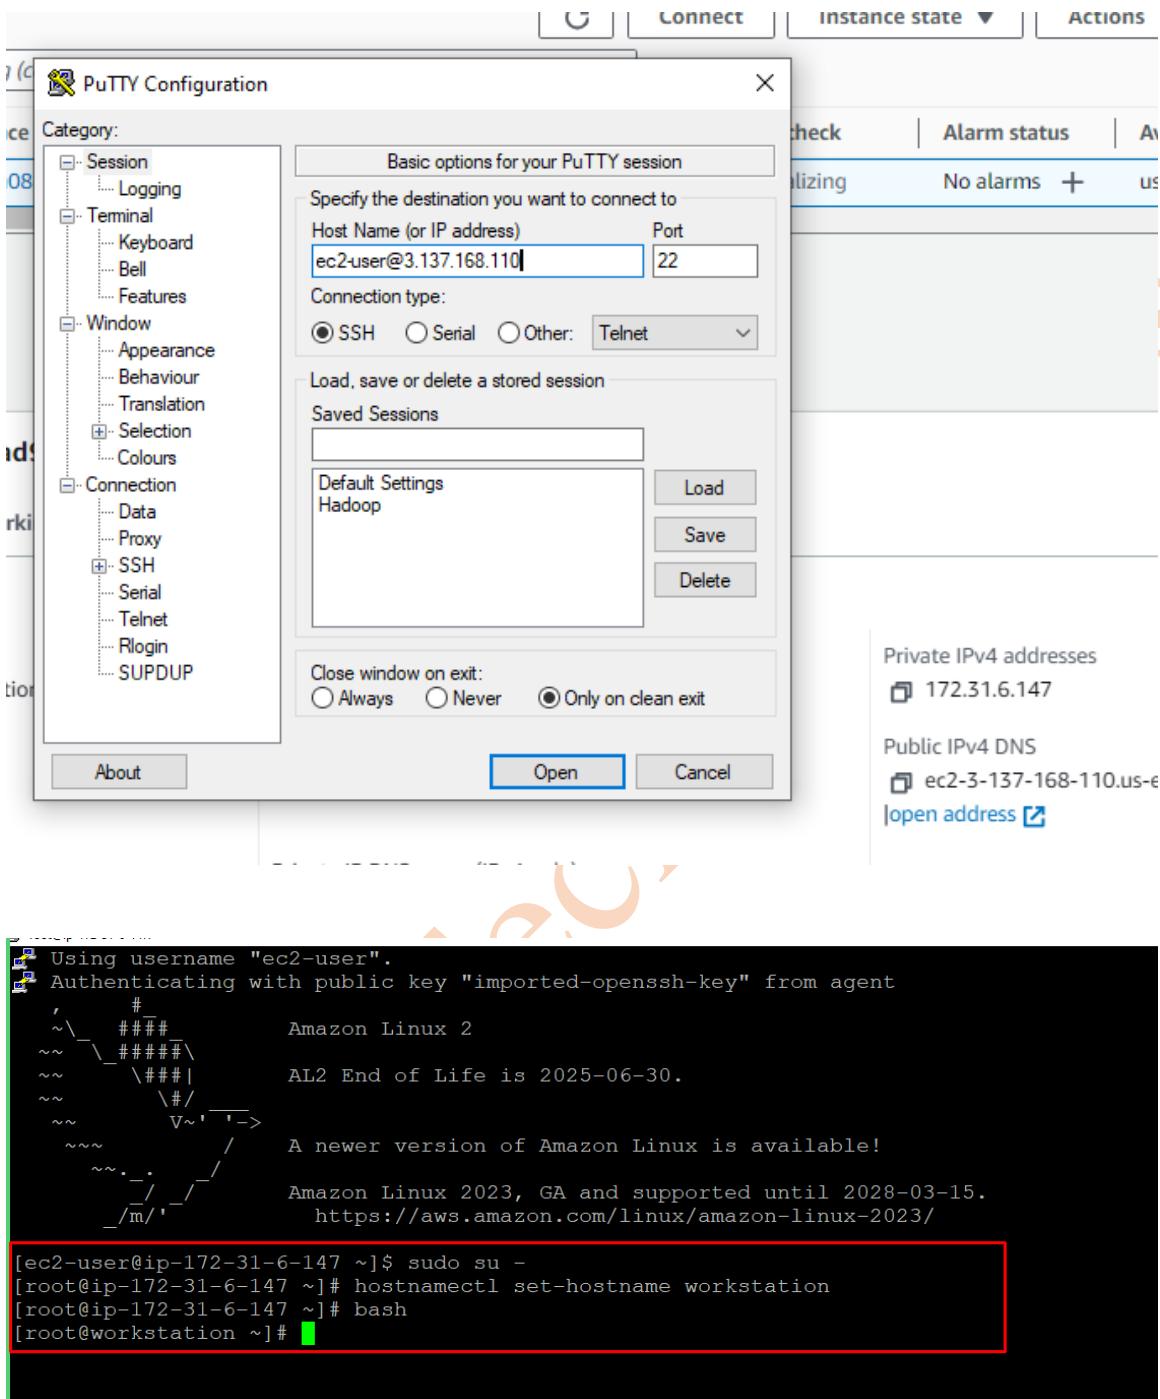


2. Login to the server and do the below

2. Login to the server and do the below







3. Installing all the required softwares to proceed with installation

1. Aws cli installation:

```
curl "https://awscli.amazonaws.com/awscli-exe-linux-x86\_64.zip" -o "awscliv2.zip"
unzip awscliv2.zip
sudo ./aws/install
```

2. Install eksctl latest binary

<https://eksctl.io/installation/>

```
# for ARM systems, set ARCH to: `arm64` , `armv6` or `armv7`
ARCH=amd64
PLATFORM=$(uname -s)_$ARCH
curl -sLO "https://github.com/eksctl-io/eksctl/releases/latest/download/eksctl\_\$PLATFORM.tar.gz"
# (Optional) Verify checksum
curl -sL "https://github.com/eksctl-io/eksctl/releases/latest/download/eksctl\_checksums.txt"
| grep $PLATFORM | sha256sum --check
tar -xzf eksctl_$PLATFORM.tar.gz -C /tmp && /usr/bin/rm -f eksctl_$PLATFORM.tar.gz
sudo mv /tmp/eksctl /usr/local/bin
echo 'export PATH=$HOME/bin:$PATH' >> ~/.bashrc
```

3. Install latest of needed version of kubectl command

<https://docs.aws.amazon.com/eks/latest/userguide/install-kubectl.html>

```
curl -O https://s3.us-west-2.amazonaws.com/amazon-eks/1.28.3/2023-11-14/bin/linux/amd64/kubectl
chmod +x ./kubectl
mkdir -p $HOME/bin && cp ./kubectl $HOME/bin/kubectl && export
PATH=$HOME/bin:$PATH
echo 'export PATH=$HOME/bin:$PATH' >> ~/.bashrc
kubectl version --client
```

```
[root@workstation ~]# aws --version
aws-cli/2.15.0 Python/3.11.6 Linux/5.10.201-191.748.amzn2.x86_64 exe/x86_64.amzn.2 prompt/off
[root@workstation ~]# eksctl version
0.165.0
[root@workstation ~]# kubectl version --client
Client Version: v1.28.3-eks-e71965b
Kustomize Version: v5.0.4-0.20230601165947-6ce0bf390ce3
[root@workstation ~]#
```

4. Configuring aws cli

The screenshot shows the AWS IAM 'Users' page. On the left, there's a sidebar with navigation links like 'Dashboard', 'Access management', 'Access reports', and 'Identity and Access Management (IAM)'. The main area displays a table of users with columns for 'User name', 'Path', 'Group', 'Last activity', 'MFA', 'Password age', 'Console last sign-in', and 'Access key ID'. A search bar at the top allows filtering by user name. At the top right, there are 'Create user' (orange), 'Delete' (grey), and other buttons. A red box highlights the 'Create user' button.

This screenshot shows the 'Specify user details' step of the IAM user creation wizard. It includes fields for 'User name' (set to 'eks-admin'), 'Provide user access to the AWS Management Console' (unchecked), and a note about generating programmatic access keys. The 'Next Step' button is highlighted with a red box.

This screenshot shows the 'Set permissions' step. It features options for 'Add user to group', 'Copy permissions', and 'Attach policies directly'. The 'Attach policies directly' option is selected and highlighted with a red box. Below, a list of 'Permissions policies' is shown, with one policy, 'AdministratorAccess', being selected and highlighted with a red box. The 'Next Step' button is also highlighted with a red box.

Set permissions boundary - optional

AWS managed	0

Cancel Previous **Next**

User name eks-admin	Console password type None	Require password reset No
------------------------	-------------------------------	------------------------------

Permissions summary

Name	Type	Used as
AdministratorAccess	AWS managed - job function	Permissions policy

Tags - optional
Tags are key-value pairs you can add to AWS resources to help identify, organize, or search for resources. Choose any tags you want to associate with this user.

No tags associated with the resource.

Add new tag
You can add up to 50 more tags.

Cancel Previous **Create user**

Now let's generate the credentials (access key,token)

User created successfully
You can view and download the user's password and email instructions for signing in to the AWS Management Console.

IAM > Users

Users (6) Info
An IAM user is an identity with long-term credentials that is used to interact with AWS in an account.

Search: eks-admin

User name	Path	Group	Last activity	MFA	Password age	Console last sign-in
eks-admin	/	0	-	-	-	-

View user Create user

Screenshot of the AWS IAM Access Management console showing the creation of a new access key.

The top navigation bar includes tabs: Permissions, Groups, Tags, **Security credentials** (highlighted with a red box), and Access Advisor.

Console sign-in section shows a "Console sign-in link" (https://bimalk-aws-v2.signin.aws.amazon.com/console) and a "Console password" status (Not enabled). A "Enable console access" button is present.

Multi-factor authentication (MFA) (0) section indicates "No MFA devices. Assign an MFA device to improve the security of your AWS environment." An "Assign MFA device" button is available.

Access keys (0) section shows a note: "Use access keys to send programmatic calls to AWS from the AWS CLI, AWS Tools for PowerShell, AWS SDKs, or direct AWS API calls. You can have a maximum of two access keys (active or inactive) at a time." A "Create access key" button is highlighted with a red box.

Use case section:

- Command Line Interface (CLI)** (selected, highlighted with a red box): "You plan to use this access key to enable the AWS CLI to access your AWS account."
- Ctrl local code**: "You plan to use this access key to enable application code in a local development environment to access your AWS account."
- Application running on an AWS compute service**: "You plan to use this access key to enable application code running on an AWS compute service like Amazon EC2, Amazon ECS, or AWS Lambda to access your AWS account."
- Third-party service**: "You plan to use this access key to enable access for a third-party application or service that ..."
- Other**: "Your use case is not listed here."

Alternatives recommended (warning icon):

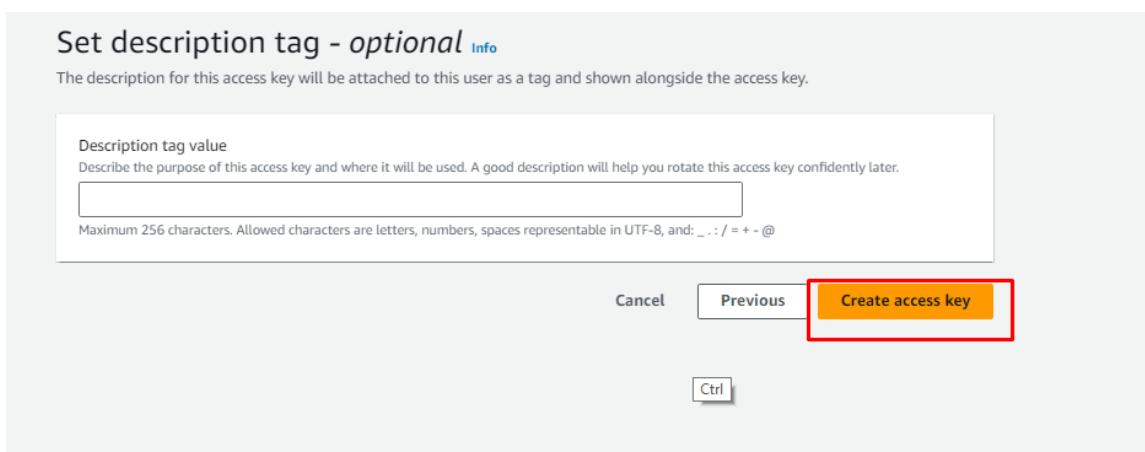
- Use [AWS CloudShell](#), a browser-based CLI, to run commands. [Learn more](#)
- Use the [AWS CLI V2](#) and enable authentication through a user in IAM Identity Center. [Learn more](#)

Confirmation section:

I understand the above recommendation and want to proceed to create an access key.

Cancel

Next



Access key	Secret access key
AKIAYSCA...	***** Show

Access key best practices

- Never store your access key in plain text, in a code repository, or in code.
- Disable or delete access key when no longer needed.
- Enable least-privilege permissions.
- Rotate access keys regularly.

For more details about managing access keys, see the [best practices for managing AWS access keys](#).

Now you can go to workstation & do the aws configure as mentioned below

```
Rustamize version: v0.0.4-0.20230601165947-6ce0b1590de3
[root@workstation ~]# aws configure
AWS Access Key ID [None]: AKIAYSCA...
AWS Secret Access Key [None]: I5peYBi6TW90...
Default region name [None]: us-east-2
Default output format [None]: json
[root@workstation ~]#
[root@workstation ~]# aws s3 ls
2023-10-16 04:31:16 destination-oregon1
2023-10-16 04:33:19 source-n.virginia
[root@workstation ~]#
```

Finally we can see that we have successfully prepared our workstation with all necessary software's required.

3 Creating cluster with eksctl

Creating cluster:

```
eksctl create cluster --name dvsekscluster --region us-east-2 --instance-types t2.medium --managed --vpc-cidr 192.168.100.0/24 --node-private-networking --version 1.27 --without-nodegroup
```

```
[root@workstation ~]# eksctl create cluster --name dvsekscluster --region us-east-2 --instance-types t2.medium --managed --vpc-cidr 192.168.100.0/24 --node-private-networking --version 1.27 --without-nodegroup
2023-12-23 07:31:50 [ ] eksctl version 0.166.0
2023-12-23 07:31:50 [ ] using region us-east-2
2023-12-23 07:31:50 [ ] setting availability zones to [us-east-2a us-east-2b us-east-2c]
2023-12-23 07:31:50 [ ] subnets for us-east-2a - public:192.168.100.0/27 private:192.168.100.96/27
2023-12-23 07:31:50 [ ] subnets for us-east-2b - public:192.168.100.32/27 private:192.168.100.128/27
2023-12-23 07:31:50 [ ] subnets for us-east-2c - public:192.168.100.64/27 private:192.168.100.160/27
2023-12-23 07:31:50 [ ] using Kubernetes version 1.27
2023-12-23 07:31:50 [ ] creating EKS cluster "dvsekscluster" in "us-east-2" region with
2023-12-23 07:31:50 [ ] if you encounter any issues, check CloudFormation console or try 'eksctl utils
describe-stacks --region=us-east-2 --cluster=dvsekscluster'
2023-12-23 07:31:50 [ ] Kubernetes API endpoint access will use default of {publicAccess=true, private
Access=false} for cluster "dvsekscluster" in "us-east-2"
2023-12-23 07:31:50 [ ] CloudWatch logging will not be enabled for cluster "dvsekscluster" in "us-east
-2"
2023-12-23 07:31:50 [ ] you can enable it with 'eksctl utils update-cluster-logging --enable-types={SP
ECIFY-YOUR-LOG-TYPES-HERE (e.g. all)} --region=us-east-2 --cluster=dvsekscluster'
2023-12-23 07:31:50 [ ]
2 sequential tasks: { create cluster control plane "dvsekscluster", wait for control plane to become re
ady}
```

Note: Cluster creation takes 10 mins, please be patient enough and wait for it to come up.

```
2023-12-23 07:36:51 [ ] waiting for CloudFormation stack "eksctl-dvsekscluster-cluster"
2023-12-23 07:37:51 [ ] waiting for CloudFormation stack "eksctl-dvsekscluster-cluster"
2023-12-23 07:38:51 [ ] waiting for CloudFormation stack "eksctl-dvsekscluster-cluster"
2023-12-23 07:39:51 [ ] waiting for CloudFormation stack "eksctl-dvsekscluster-cluster"
2023-12-23 07:40:51 [ ] waiting for CloudFormation stack "eksctl-dvsekscluster-cluster"

2023-12-23 07:42:52 [ ] waiting for the control plane to become ready
2023-12-23 07:42:53 [v] saved kubeconfig as "/root/.kube/config"
2023-12-23 07:42:53 [ ] no tasks
2023-12-23 07:42:53 [v] all EKS cluster resources for "dvsekscluster" have been created
2023-12-23 07:42:54 [ ] kubectl command should work with "/root/.kube/config", try 'kubectl get nodes'
2023-12-23 07:42:54 [v] EKS cluster "dvsekscluster" in "us-east-2" region is ready
```

```
[root@workstation ~]# aws eks list-clusters --region us-east-2
{
    "clusters": [
        "dvsekscluster"
    ]
}
[root@workstation ~]#
```

Now we can see that our cluster is up & running.

Creating custom network:

Update the kubeconfig file with below command to login to the cluster

```
aws eks update-kubeconfig --region us-east-2 --name "dvsekscluster"
```

```
[root@workstation ~]# aws eks update-kubeconfig --region us-east-2 --name "dvsekscluster"
Added new context arn:aws:eks:us-east-2:439292780570:cluster/dvsekscluster to /root/.kube/config
[root@workstation ~]#
```

kubectl delete daemonset -n kube-system aws-node

kubectl create -f

<https://raw.githubusercontent.com/projectcalico/calico/v3.25.1/manifests/tigera-operator.yaml>

```
[root@workstation ~]# kubectl delete daemonset -n kube-system aws-node
daemonset.apps "aws-node" deleted
[root@workstation ~]# kubectl create -f https://raw.githubusercontent.com/projectcalico/calico/v3.25.1/
manifests/tigera-operator.yaml
namespace/tigera-operator created
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
```

kubectl create -f - <<EOF

kind: Installation

apiVersion: operator.tigera.io/v1

metadata:

name: default

spec:

kubernetesProvider: EKS

cni:

type: Calico

calicoNetwork:

bgp: Disabled

EOF

```
[root@workstation ~]# kubectl create -f - <<EOF
> kind: Installation
> apiVersion: operator.tigera.io/v1
> metadata:
>   name: default
> spec:
>   kubernetesProvider: EKS
>   cni:
>     type: Calico
>   calicoNetwork:
>     bgp: Disabled
> EOF
installation.operator.tigera.io/default created
[root@workstation ~]#
```

Creating nodegroup:

```
eksctl create nodegroup --cluster "dvsekscluster" --name nodegroup1 --nodes 2 --node-type t2.medium --node-private-networking --managed --max-pods-per-node 200
```

```
[root@workstation ~]#
[root@workstation ~]# eksctl create nodegroup --cluster "dvsekscluster" --name nodegroup1 --nodes 2 --node-type t2.medium --node-private-networking --managed --max-pods-per-node 200
2023-12-23 07:51:24 [ ] will use version 1.27 for new nodegroup(s) based on control plane version
2023-12-23 07:51:25 [ ] nodegroup "nodegroup1" will use "" (AmazonLinux2/1.27)
2023-12-23 07:51:25 [ ] 1 nodegroup (nodegroup1) was included (based on the include/exclude rules)
2023-12-23 07:51:25 [ ] will create a CloudFormation stack for each of 1 managed nodegroups in cluster "dvsekscluster"
2023-12-23 07:51:25 [ ]
2 sequential tasks: { fix cluster compatibility, 1 task: { 1 task: { create managed nodegroup "nodegroup1" } }
}
2023-12-23 07:51:25 [ ] checking cluster stack for missing resources
2023-12-23 07:51:25 [ ] cluster stack has all required resources
2023-12-23 07:51:25 [ ] building managed nodegroup stack "eksctl-dvsekscluster-nodegroup-nodegroup1"
2023-12-23 07:51:25 [ ] deploying stack "eksctl-dvsekscluster-nodegroup-nodegroup1"
2023-12-23 07:51:25 [ ] waiting for CloudFormation stack "eksctl-dvsekscluster-nodegroup-nodegroup1"
```

Here we are going to create the worker nodes which we call it as node group in eks cluster with two nodes of hardware type t2.medium.

Final verification:

```
[root@workstation ~]# kubectl get nodes
NAME                               STATUS    ROLES      AGE      VERSION
ip-192-168-100-110.us-east-2.compute.internal  Ready    <none>    50m      v1.27.7-eks-e71965b
ip-192-168-100-173.us-east-2.compute.internal  Ready    <none>    50m      v1.27.7-eks-e71965b
[root@workstation ~]# kubectl get ns
NAME          STATUS  AGE
calico-system  Active  50m
default        Active  67m
kube-node-lease Active  67m
kube-public    Active  67m
kube-system    Active  67m
tigera-operator Active  58m
[root@workstation ~]#
```

4. Working With Cluster

Namespace Creation:

```
vi ns.yaml
```

```
apiVersion: v1
kind: Namespace
metadata:
  creationTimestamp: null
  name: myapplication
spec: {}
status: {}
```

```
[root@workstation myobjects]# vi ns.yaml
[root@workstation myobjects]# cat ns.yaml
apiVersion: v1
kind: Namespace
metadata:
  creationTimestamp: null
  name: myapplication
spec: {}
status: {}
[root@workstation myobjects]# kubectl apply -f ns.yaml
namespace/myapplication created
[root@workstation myobjects]# kubectl get ns
NAME      STATUS   AGE
calico-system  Active  60m
default     Active  77m
kube-node-lease  Active  77m
kube-public   Active  77m
kube-system   Active  77m
myapplication  Active  4s
tigera-operator  Active  69m
[root@workstation myobjects]#
```

Pod Creation:

```
vi pod.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: mypod
    name: mypod
    namespace: myapplication
spec:
  containers:
    - image: nginx
```

```
name: mypod
resources: {}
dnsPolicy: ClusterFirst
restartPolicy: Always
status: {}
```

```
[root@workstation myobjects]# cat pod.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: mypod
  name: mypod
  namespace: myapplication
spec:
  containers:
  - image: nginx
    name: mypod
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
[root@workstation myobjects]# kubectl apply -f pod.yaml -n myapplication
pod/mypod created
[root@workstation myobjects]# kubectl get pods -n myapplication
NAME      READY   STATUS    RESTARTS   AGE
mypod    1/1     Running   0          10s
[root@workstation myobjects]#
```

Deployment creation:

vi deploy.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
  name: mydeployment
  namespace: myapplication
spec:
  replicas: 1
  selector:
    matchLabels:
      app: mydeployment
  strategy: {}
  template:
    metadata:
```

```

creationTimestamp: null
labels:
  app: mydeployment
spec:
  containers:
    - image: nginx
      name: nginx
      resources: {}
status: {}

```

```

[root@workstation myobjects]# cat deploy.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
    name: mydeployment
    namespace: myapplication
spec:
  replicas: 1
  selector:
    matchLabels:
      app: mydeployment
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: mydeployment
    spec:
      containers:
        - image: nginx
          name: nginx
          resources: {}
status: {}

```

```

status: {}
[root@workstation myobjects]# kubectl apply -f deploy.yaml
deployment.apps/mydeployment created
[root@workstation myobjects]# kubectl get deploy -n myapplication
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
mydeployment   1/1     1           1           10s
[root@workstation myobjects]#

```

```

[root@workstation:~/myobjects]
[root@workstation myobjects]# kubectl get pods -n myapplication
NAME           READY   STATUS    RESTARTS   AGE
mydeployment-8496ddbc68-2zxxs  1/1     Running   0          66s
mypod          1/1     Running   0          5m34s
[root@workstation myobjects]# kubectl get pods -n myapplication -l app=mydeployment
NAME           READY   STATUS    RESTARTS   AGE
mydeployment-8496ddbc68-2zxxs  1/1     Running   0          75s
[root@workstation myobjects]#

```

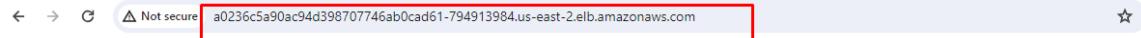
Service creation:

```
vi svc.yaml
```

```
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
  name: mydeployment
  namespace: myapplication
spec:
  ports:
  - port: 80
    protocol: TCP
    targetPort: 80
  selector:
    app: mydeployment
  type: LoadBalancer
status:
  loadBalancer: {}
```

```
[root@workstation myobjects]# [root@workstation myobjects]# cat svc.yaml
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
  name: mydeployment
  namespace: myapplication
spec:
  ports:
  - port: 80
    protocol: TCP
    targetPort: 80
  selector:
    app: mydeployment
  type: LoadBalancer
status:
  loadBalancer: {}
[root@workstation myobjects]# kubectl apply -f svc.yaml
service/mydeployment created
[root@workstation myobjects]# kubectl get svc -n myapplication
NAME      TYPE           CLUSTER-IP   EXTERNAL-IP
mydeployment   LoadBalancer   10.100.150.229   a0236c5a90ac94d398707746ab0cad61-794913984.us-east-2.elb.amazonaws.com   80:30507/TCP   11s
[root@workstation myobjects]# kubectl get endpoints -n myapplication
```

We can able to see that our deployment got exposed as service of type loadbalancer. We got out LB url which we can use it for accessing our application from browser.



Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.

If you want to see the pods behind this service then you can execute the below command

```
[root@workstation myobjects]# kubectl get endpoints -n myapplication
NAME      ENDPOINTS   AGE
mydeployment  172.16.249.67:80  3m46s
[root@workstation myobjects]# kubectl get pods -n myapplication -o wide
NAME          READY   STATUS    RESTARTS   AGE   NODE
mydeployment-8496ddbc68-2zxxs  1/1     Running   0          10m   172.16.249.67   ip-192-168-100-173.us-east-2.compute.internal
mypod        1/1     Running   0          14m   172.16.249.66   ip-192-168-100-173.us-east-2.compute.internal
[root@workstation myobjects]#
```

Increasing Replica Count:

If you want to run more than one pod then you can increase the replica count in the below yaml

```
[root@workstation myobjects]# cat deploy.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
  name: mydeployment
  namespace: myapplication
spec:
  replicas: 2
  selector:
    matchLabels:
      app: mydeployment
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: mydeployment
    spec:
      containers:
        - image: nginx
          name: nginx
          resources: {}
status: {}
```

Before changes:

```
[root@workstation myobjects]# kubectl get po -n myapplication -l app=mydeployment
NAME           READY   STATUS    RESTARTS   AGE
mydeployment-8496ddbc68-kbn65  1/1     Running   0          16h
[root@workstation myobjects]# kubectl apply -f deploy.yaml
```

We can see before we apply the yaml no.of pods for this deployment is 2.

Post changes:

```
[root@workstation myobjects]# kubectl apply -f deploy.yaml
deployment.apps/mydeployment configured
[root@workstation myobjects]# kubectl get po -n myapplication -l app=mydeployment
NAME           READY   STATUS    RESTARTS   AGE
mydeployment-8496ddbc68-chxdc  1/1     Running   0          7s
mydeployment-8496ddbc68-kbn65  1/1     Running   0          16h
[root@workstation myobjects]#
```

Overall Commands:

```
[root@workstation:~/myobjects]
[root@workstation myobjects]# kubectl get ns
NAME      STATUS  AGE
calico-system  Active  17h
default    Active  18h
kube-node-lease  Active  18h
kube-public  Active  18h
kube-system  Active  18h
myapplication  Active  16h
tigera-operator  Active  17h
[root@workstation myobjects]# kubectl get deploy -n myapplication
NAME        READY  UP-TO-DATE  AVAILABLE  AGE
mydeployment  2/2    2          2          16h
[root@workstation myobjects]# kubectl get pods -n myapplication -l app=mydeployment
NAME           READY   STATUS    RESTARTS   AGE
mydeployment-8496ddbc68-chxdc  1/1     Running   0          2m9s
mydeployment-8496ddbc68-kbn65  1/1     Running   0          16h
[root@workstation myobjects]# kubectl get svc -n myapplication
NAME        TYPE        CLUSTER-IP      EXTERNAL-IP
mydeployment  LoadBalancer  10.100.150.229  a0236c5a90ac94d398707746ab0cad61-794913984.us-east-2.elb.amazonaws.com  80:30507/TCP  16h
[root@workstation myobjects]# kubectl get endpoints -n myapplication
NAME        ENDPOINTS
mydeployment  172.16.99.1:80,172.16.99.6:80  16h
[root@workstation myobjects]#
```

5. Working with Service

Service acts like as a loadbalancer on top of all the pods of an expose deployment. We have three types of services in kubernetes. They are clusterip,nodeport & loadbalancer

1. **ClusterIP:** (This is a default service type, which helps communication between microservices internal to the cluster)

```
[root@workstation myobjects]# kubectl get deploy -n myapplication
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
mydeployment   2/2     2            2           16h
[root@workstation myobjects]# kubectl get pods -n myapplication -l app=mydeployment
NAME             READY   STATUS    RESTARTS   AGE
mydeployment-8496ddbc68-chxdc  1/1     Running   0          10m
mydeployment-8496ddbc68-kbn65  1/1     Running   0          16h
[root@workstation myobjects]# kubectl get pods -n myapplication -l app=mydeployment -o wide
NAME             READY   STATUS    RESTARTS   AGE   NOMINATED NODE   READINESS GATES
mydeployment-8496ddbc68-chxdc  1/1     Running   0          10m   172.16.99.6   ip-192-168-100-171.us-east-2.compute.internal   <none>        <none>
mydeployment-8496ddbc68-kbn65  1/1     Running   0          16h   172.16.99.1   ip-192-168-100-171.us-east-2.compute.internal   <none>        <none>
[root@workstation myobjects]# hostname -I
172.31.21.41 172.17.0.1
[root@workstation myobjects]#
```

Above screenshot give us an overview of no of pods running for a deployment and their ipaddresses. If you observe properly our host ipaddress(workstation) is not in the range of the cluster ipaddress hence we need an intermediate container which help us to test the application. Hence we are going to use busbox container.

```
vi busybox.yaml
apiVersion: v1
kind: Pod
metadata:
  name: busybox1
  labels:
    app: busybox1
  namespace: myapplication
spec:
  containers:
    - image: radial/busyboxplus:curl
      command:
        - sleep
        - "3600"
  imagePullPolicy: IfNotPresent
  name: busybox
  restartPolicy: Always
```

```

[root@workstation myobjects]# cat busybox.yaml
apiVersion: v1
kind: Pod
metadata:
  name: busybox1
  labels:
    app: busybox1
  namespace: myapplication
spec:
  containers:
  - image: radial/busyboxplus:curl
    command:
    - sleep
    - "3600"
    imagePullPolicy: IfNotPresent
    name: busybox
    restartPolicy: Always
[root@workstation myobjects]# kubectl apply -f busybox.yaml
pod/busybox1 created
[root@workstation myobjects]# kubectl get pods -n myapplication -l app=mydeployment -o wide
NAME           READY   STATUS    RESTARTS   AGE      IP          NODE
mydeployment-8496ddbc68-chxdc  1/1     Running   0          38m     172.16.99.6   ip-192-168-100-171.us-east-2.compute.internal
mydeployment-8496ddbc68-kbn65  1/1     Running   0          17h     172.16.99.1   ip-192-168-100-171.us-east-2.compute.internal
[root@workstation myobjects]# kubectl get pods -n myapplication -l app=busybox
NAME           READY   STATUS    RESTARTS   AGE
busybox1       1/1     Running   0          22s
[root@workstation myobjects]#

```

Let's login to the busybox pod and check the application status of our deployment pod

kubectl exec -it busybox1 -n myapplication /bin/sh

```

[root@workstation myobjects]# kubectl get pods -n myapplication -l app=mydeployment -o wide
NAME           READY   STATUS    RESTARTS   AGE      IP          NODE
mydeployment-8496ddbc68-chxdc  1/1     Running   0          38m     172.16.99.6   ip-192-168-100-171.us-east-2.compute.internal
mydeployment-8496ddbc68-kbn65  1/1     Running   0          17h     172.16.99.1   ip-192-168-100-171.us-east-2.compute.internal
[root@workstation myobjects]# kubectl get pods -n myapplication -l app=busybox
NAME           READY   STATUS    RESTARTS   AGE
busybox1       1/1     Running   0          22s
[root@workstation myobjects]# kubectl exec -it busybox1 -n myapplication /bin/sh
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
/bin/sh: shopt: not found
[ root@busybox1:/ ]$ curl http://172.16.99.6
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and working. Further configuration is required.</p>

```

Let's test the application with the service of type clusterip using below.

```

vi svc-cip.yaml
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null

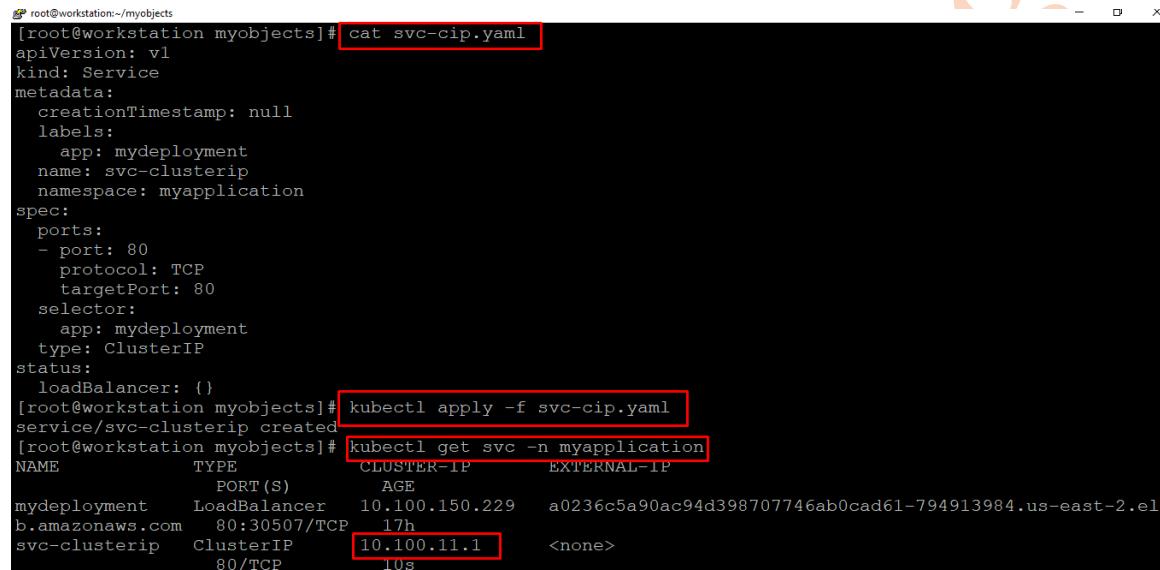
```

DVS Technologies, Opp Home Town, Beside Biryani Zone, Marathahalli, Bangalore Phone: 9632558585 Mobile: 8892499499 Mail : dvs.training@gmail.com Web: www.dvstechnologies.in

```

labels:
  app: mydeployment
  name: svc-clusterip
  namespace: myapplication
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
  selector:
    app: mydeployment
  type: ClusterIP

```



```

root@workstation:~/myobjects
[root@workstation myobjects]# cat svc-cip.yaml
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
    name: svc-clusterip
    namespace: myapplication
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
  selector:
    app: mydeployment
  type: ClusterIP
status:
  loadBalancer: {}
[root@workstation myobjects]# kubectl apply -f svc-cip.yaml
service/svc-clusterip created
[root@workstation myobjects]# kubectl get svc -n myapplication
NAME         TYPE           CLUSTER-IP      EXTERNAL-IP
mydeployment   LoadBalancer   10.100.150.229   a0236c5a90ac94d398707746ab0cad61-794913984.us-east-2.el
b.amazonaws.com  80:30507/TCP  17h
svc-clusterip   ClusterIP     10.100.11.1    <none>
               80/TCP        10s

```

Verification:

```

service/svc-clusterip created
[root@workstation myobjects]# kubectl get svc -n myapplication
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP
mydeployment   LoadBalancer  10.100.150.229  a0236c5a90ac94d398707746ab0cad61-794913984.us-east-2.el
b.amazonaws.com  80:30507/TCP  17h
svc-clusterip   ClusterIP   10.100.11.1    <none>
              80/TCP     10s
[root@workstation myobjects]# kubectl exec -it busybox1 -n myapplication /bin/sh
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
/bin/sh: shopt: not found
[ root@busybox1:/ ]$ curl http://10.100.11.1
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

```

From the above we can conclude that we can able to hit the application via service of type clusterip

2. Nodeport:

```

mydeployment  3/3    3    3    24h
[ec2-user@master ~]$ kubectl expose deploy dvsbatch3 --name=nodeport-service --target-port=80 --port 8080 --type NodePort
service/nodeport-service exposed
[ec2-user@master ~]$ kubectl get svc
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)        AGE
dvsbatch3-service  ClusterIP  10.110.222.174  <none>       8080/TCP     13m
kubernetes     ClusterIP  10.96.0.1    <none>       443/TCP      24h
myservice       ClusterIP  10.108.174.113   <none>       80/TCP       24h
nodeport-service  NodePort   10.98.235.212   <none>       8080:30964/TCP  4s
[ec2-user@master ~]$ 

```

```

vi svc-np.yaml
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
    name: svc-nodeport
    namespace: myapplication
spec:
  ports:
  - port: 80
    protocol: TCP
    targetPort: 80
  selector:
    app: mydeployment

```

type: NodePort

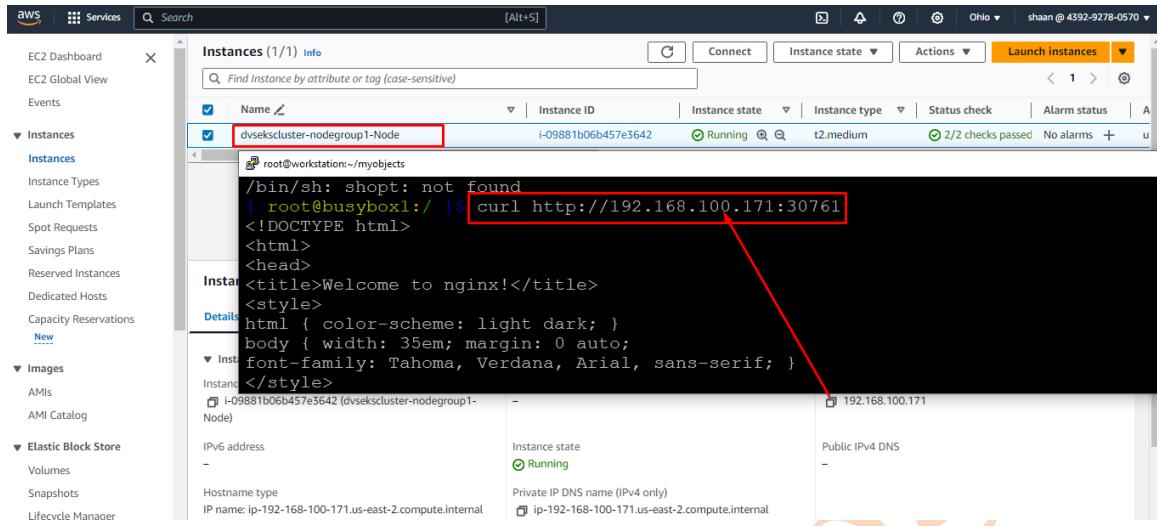
```
[root@workstation myobjects]# cat svc-np.yaml
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
  name: svc-nodeport
  namespace: myapplication
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
    selector:
      app: mydeployment
    type: NodePort
status:
  loadBalancer: {}

[root@workstation myobjects]# kubectl apply -f svc-np.yaml
service/svc-nodeport created
[root@workstation myobjects]# kubectl get svc -n myapplication
NAME           TYPE        CLUSTER-IP   EXTERNAL-IP
mydeployment   LoadBalancer 10.100.150.229 a0236c5a90ac94d398707746ab0cad61-794913984.us-east-2.elb.amazonaws.com 80:30507/TCP 17h
svc-clusterip  ClusterIP   10.100.11.1  <none>
                80/TCP     4m19s
svc-nodeport   NodePort   10.100.112.18 <none>
                80:30761/TCP 8s
[root@workstation myobjects]#
```

Verification:

```
[root@workstation:~/myobjects]
[root@workstation myobjects]# kubectl get svc -n myapplication
NAME           TYPE        CLUSTER-IP   EXTERNAL-IP
mydeployment   LoadBalancer 10.100.150.229 a0236c5a90ac94d398707746ab0cad61-794913984.us-east-2.elb.amazonaws.com 80:30507/TCP 17h
svc-clusterip  ClusterIP   10.100.11.1  <none>
                80/TCP     4m19s
svc-nodeport   NodePort   10.100.112.18 <none>
                80:30761/TCP 8s
[root@workstation myobjects]# kubectl exec -it busybox -n myapplication /bin/sh
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
/bin/sh: shopt: not found
[ root@busybox:/ ]$ curl http://192.168.100.171:30761
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>
</body>
</html>
```

Here 192.168.100.171 is the ipaddress of our workernode. You can take any worker node ipaddress to test your node port.



3. Configuring Loadbalancer type:

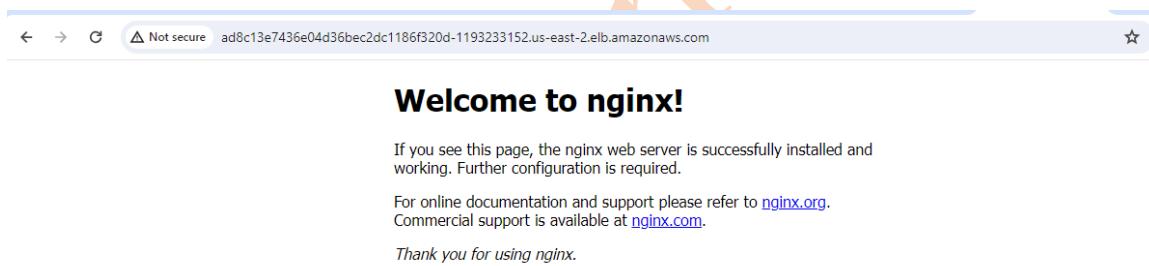
```

vi svc-lb.yaml
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
    name: svc-nodeport
    namespace: myapplication
spec:
  ports:
  - port: 80
    protocol: TCP
    targetPort: 80
  selector:
    app: mydeployment
  type: LoadBalancer
status:
  loadBalancer: {}

```

```
[root@workstation myobjects]# kubectl get pods -n myapplication -l app=mydeployment
NAME           READY   STATUS    RESTARTS   AGE
mydeployment-8496ddbc68-cdlbg   1/1     Running   0          54s
mydeployment-8496ddbc68-srbg4   1/1     Running   0          54s
[root@workstation myobjects]# cat svc-lb.yaml
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
  name: svc-nodeport
  namespace: myapplication
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
  selector:
    app: mydeployment
    type: LoadBalancer
  status:
    loadBalancer: {}
[root@workstation myobjects]# kubectl apply -f svc-lb.yaml
service/svc-nodeport created
[root@workstation myobjects]# kubectl get svc -n myapplication
NAME         TYPE        CLUSTER-IP      EXTERNAL-IP
PORT(S)      AGE
svc-nodeport LoadBalancer 10.100.221.64  ad8c13e7436e04d36bec2dc1186f320d-1193233152.us-east-2.elb.amazonaws.com  80:31748/TCP  11s
[root@workstation myobjects]#
```

From the above we got our loadbalancer url, just hit above url in the browser



Finally we got our application url and its up & running.

6. Scheduling

1. Labeling the nodes

```
ip-192-168-100-171.us-east-2.compute.internal Ready <none> 3h55m v1.27.7-eks-e71965b
[root@workstation myobjects]# kubectl get nodes
NAME STATUS ROLES AGE VERSION
ip-192-168-100-118.us-east-2.compute.internal Ready <none> 10m v1.27.7-eks-e71965b
ip-192-168-100-171.us-east-2.compute.internal Ready <none> 6h7m v1.27.7-eks-e71965b
[root@workstation myobjects]# kubectl get nodes --show-labels
NAME STATUS ROLES AGE VERSION LABELS
ip-192-168-100-118.us-east-2.compute.internal Ready <none> 10m v1.27.7-eks-e71965b alpha.e
sctl.io/cluster-name=dvsekscluster,alpha.eksctl.io/nodegroup-name=nodegroup1,beta.kubernetes.io/arch=am
d64,beta.kubernetes.io/instance-type=t2.medium,beta.kubernetes.io/os=linux,eks.amazonaws.com/capacityT
pe=ON_DEMAND,eks.amazonaws.com/nodegroup-image=ami-02d77f9fdc5a964ea,eks.amazonaws.com/nodegroup=nod
eroup1,eks.amazonaws.com/sourceLaunchTemplateId=lt-0f742ac99001ba24f,eks.amazonaws.com/sourceLaunchTempl
ateVersion=1,failure-domain.beta.kubernetes.io/region=us-east-2,failure-domain.beta.kubernetes.io/zone=
s-east-2a,k8s.io/cloud-provider-aws=c9047fecf124ac306f038a64fec27c52,kubernetes.io/arch=amd64,kuberne
tes.io/hostname=ip-192-168-100-118.us-east-2.compute.internal,kubernetes.io/os=linux,node.kubernetes.io/
instance-type=t2.medium,topology.kubernetes.io/region=us-east-2,topology.kubernetes.io/zone=us-east-2a
ip-192-168-100-171.us-east-2.compute.internal Ready <none> 6h7m v1.27.7-eks-e71965b alpha.e
sctl.io/cluster-name=dvsekscluster,alpha.eksctl.io/nodegroup-name=nodegroup1,beta.kubernetes.io/arch=am
d64,beta.kubernetes.io/instance-type=t2.medium,beta.kubernetes.io/os=linux,eks.amazonaws.com/capacityT
pe=ON_DEMAND,eks.amazonaws.com/nodegroup-image=ami-02d77f9fdc5a964ea,eks.amazonaws.com/nodegroup=nod
eroup1,eks.amazonaws.com/sourceLaunchTemplateId=lt-0f742ac99001ba24f,eks.amazonaws.com/sourceLaunchTempl
ateVersion=1,failure-domain.beta.kubernetes.io/region=us-east-2,failure-domain.beta.kubernetes.io/zone=
s-east-2c,k8s.io/cloud-provider-aws=c9047fecf124ac306f038a64fec27c52,kubernetes.io/arch=amd64,kuberne
tes.io/hostname=ip-192-168-100-171.us-east-2.compute.internal,kubernetes.io/os=linux,node.kubernetes.io/
instance-type=t2.medium,topology.kubernetes.io/region=us-east-2,topology.kubernetes.io/zone=us-east-2c
[root@workstation myobjects]#
```

Adding label to a node:

```
^[[A
No resources found
[root@workstation myobjects]# kubectl get nodes -l env=dev
No resources found
[root@workstation myobjects]# kubectl label node ip-192-168-100-171.us-east-2.compute.internal env=dev
node/ip-192-168-100-171.us-east-2.compute.internal labeled
[root@workstation myobjects]# kubectl get nodes -l env=dev
NAME STATUS ROLES AGE VERSION
ip-192-168-100-171.us-east-2.compute.internal Ready <none> 4h52m v1.27.7-eks-e71965b
[root@workstation myobjects]#
```

Remove label from a node:

```
[root@workstation myobjects]# kubectl get nodes -l env=dev
NAME STATUS ROLES AGE VERSION
ip-192-168-100-171.us-east-2.compute.internal Ready <none> 4h54m v1.27.7-eks-e71965b
[root@workstation myobjects]# kubectl label node ip-192-168-100-171.us-east-2.compute.internal env-
node/ip-192-168-100-171.us-east-2.compute.internal unlabeled
[root@workstation myobjects]# kubectl get nodes -l env=dev
No resources found
[root@workstation myobjects]#
```

2 . Nodename:

If you want to schedule pods in only one node then you can use the below

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
```

```

name: mydeployment
namespace: myapplication
spec:
replicas: 2
selector:
matchLabels:
app: mydeployment
strategy: {}
template:
metadata:
creationTimestamp: null
labels:
app: mydeployment
spec:
containers:
- image: nginx
name: nginx
resources: {}
nodeName: ip-192-168-100-118.us-east-2.compute.internal
status: {}

```

```

[root@workstation myobjects]# kubectl apply -f deploy.yaml
deployment.apps/mydeployment configured
[root@workstation myobjects]# kubectl get pods -n myapplication -l app=mydeployment -o wide
NAME           READY   STATUS    RESTARTS   AGE   IP          NODE
mydeployment-868b9c8fd6-h6c4m  1/1     Running   0          2m2s  172.16.71.67  ip-192-168-100-118.us-east-2.compute.internal
mydeployment-868b9c8fd6-pottz  1/1     Running   0          2m7s  172.16.71.66  ip-192-168-100-118.us-east-2.compute.internal
[root@workstation myobjects]#

```

Let's increase the replica count & see if the pods get schedule in different nodes or same node.

```

root@workstation:~/myobjects
[root@workstation myobjects]# cat deploy.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
    name: mydeployment
    namespace: myapplication
spec:
  replicas: 4
  selector:
    matchLabels:
      app: mydeployment
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: mydeployment
    spec:
      containers:
        - image: nginx
          name: nginx
          resources: {}
      nodeName: ip-192-168-100-118.us-east-2.compute.internal
status: {}
[root@workstation myobjects]# kubectl apply -f deploy.yaml
deployment.apps/mydeployment configured
[root@workstation myobjects]#

```

```

root@workstation:~/myobjects
[root@workstation myobjects]# kubectl get pods -n myapplication -l app=mydeployment -o wide
NAME                         READY   STATUS    RESTARTS   AGE   IP           NODE
mydeployment-868b9c8fd6-948p2  1/1    Running   0          17s   172.16.71.68  ip-192-168-100-118.us-east-2.compute.internal
mydeployment-868b9c8fd6-h6c4m   1/1    Running   0          3m24s  172.16.71.67  ip-192-168-100-118.us-east-2.compute.internal
mydeployment-868b9c8fd6-kvqlh   1/1    Running   0          17s   172.16.71.69  ip-192-168-100-118.us-east-2.compute.internal
mydeployment-868b9c8fd6-pcttz   1/1    Running   0          3m29s  172.16.71.66  ip-192-168-100-118.us-east-2.compute.internal
[root@workstation myobjects]#

```

From the above output we can conclude that all the 4 pods are running in the same node.

2. NodeSelector

```

us-east-2.compute.internal [root@workstation myobjects]# kubectl get nodes
[root@workstation myobjects]# kubectl get nodes
NAME                     STATUS   ROLES   AGE     VERSION
ip-192-168-100-118.us-east-2.compute.internal Ready   <none>  22m    v1.27.7-eks-e71965b
ip-192-168-100-171.us-east-2.compute.internal Ready   <none>  6h19m   v1.27.7-eks-e71965b
[root@workstation myobjects]# kubectl get nodes -l env=dev
No resources found
[root@workstation myobjects]# kubectl label node ip-192-168-100-171.us-east-2.compute.internal env=dev
node/ip-192-168-100-171.us-east-2.compute.internal labeled
[root@workstation myobjects]# kubectl get nodes -l env=dev
NAME                     STATUS   ROLES   AGE     VERSION
ip-192-168-100-171.us-east-2.compute.internal Ready   <none>  6h20m   v1.27.7-eks-e71965b
[root@workstation myobjects]#

```

Let's create the deployment with nodeselector.

```
[root@workstation myobjects]# cat deploy.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
    name: mydeployment
    namespace: myapplication
spec:
  replicas: 2
  selector:
    matchLabels:
      app: mydeployment
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: mydeployment
    spec:
      containers:
        - image: nginx
          name: nginx
          resources: {}
      nodeSelector:
        env: dev
  status: {}
```

DVS Technologies

```
[root@workstation myobjects]# cat deploy.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
  name: mydeployment
  namespace: myapplication
spec:
  replicas: 2
  selector:
    matchLabels:
      app: mydeployment
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: mydeployment
    spec:
      containers:
        - image: nginx
          name: nginx
          resources: {}
      nodeSelector:
        env: dev
status: {}
[root@workstation myobjects]# kubectl apply -f deploy.yaml
deployment.apps/mydeployment configured
[root@workstation myobjects]# kubectl get pods -n myapplication -l app=mydeployment -o wide
```

From the above we can see that our deployment is going to run on the nodes who is having label as "env=dev".

```
[root@workstation myobjects]# kubectl get nodes -l env=dev
NAME           STATUS   ROLES   AGE     VERSION
ip-192-168-100-171.us-east-2.compute.internal Ready   <none>  6h29m   v1.27.7-eks-e71965b
[root@workstation myobjects]# kubectl get pods -n myapplication -l app=mydeployment -o wide
NAME           READY   STATUS   RESTARTS   AGE     IP           NODE
mydeployment-665d88877-lp27x  1/1    Running   0          2m16s   172.16.99.16  ip-192-168-100-171.us-east-2.compute.internal
mydeployment-665d88877-mhzlk  1/1    Running   0          2m14s   172.16.99.17  ip-192-168-100-171.us-east-2.compute.internal
[root@workstation myobjects]#
```

Now let's give label as env=dev to other node & then increase the replica count to 4 & check where the new pods going to get schedule.

```
[root@workstation myobjects]# kubectl get nodes -l env=dev
NAME           STATUS   ROLES   AGE     VERSION
ip-192-168-100-118.us-east-2.compute.internal Ready   <none>  39m    v1.27.7-eks-e71965b
ip-192-168-100-171.us-east-2.compute.internal Ready   <none>  6h35m   v1.27.7-eks-e71965b
[root@workstation myobjects]# kubectl get nodes -l env=dev
NAME           STATUS   ROLES   AGE     VERSION
ip-192-168-100-171.us-east-2.compute.internal Ready   <none>  6h36m   v1.27.7-eks-e71965b
[root@workstation myobjects]# kubectl label node ip-192-168-100-118.us-east-2.compute.internal env=dev
node/ip-192-168-100-118.us-east-2.compute.internal labeled
[root@workstation myobjects]# kubectl get nodes -l env=dev
NAME           STATUS   ROLES   AGE     VERSION
ip-192-168-100-118.us-east-2.compute.internal Ready   <none>  39m    v1.27.7-eks-e71965b
ip-192-168-100-171.us-east-2.compute.internal Ready   <none>  6h36m   v1.27.7-eks-e71965b
[root@workstation myobjects]#
```

```
[root@workstation myobjects]# cat deploy.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
  name: mydeployment
  namespace: myapplication
spec:
  replicas: 4
  selector:
    matchLabels:
      app: mydeployment
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: mydeployment
    spec:
      containers:
        - image: nginx
          name: nginx
          resources: {}
  nodeSelector:
    env: dev
```

```
[root@workstation myobjects]# kubectl get pods -n myapplication -l app=mydeployment -o wide
NAME           READY   STATUS    RESTARTS   AGE   IP           NODE
mydeployment-665d88877-1p27x  1/1    Running   0          10m   172.16.99.16   ip-192-168-100-171.us-east-2.compute.internal
mydeployment-665d88877-mhzlk  1/1    Running   0          10m   172.16.99.17   ip-192-168-100-171.us-east-2.compute.internal
[root@workstation myobjects]# kubectl apply -f deploy.yaml
deployment.apps/mydeployment configured
[root@workstation myobjects]# kubectl get pods -n myapplication -l app=mydeployment -o wide
NAME           READY   STATUS    RESTARTS   AGE   IP           NODE
mydeployment-665d88877-bcpwh  1/1    Running   0          3s    172.16.71.71   ip-192-168-100-118.us-east-2.compute.internal
mydeployment-665d88877-1p27x  1/1    Running   0          10m   172.16.99.16   ip-192-168-100-171.us-east-2.compute.internal
mydeployment-665d88877-mhzlk  1/1    Running   0          10m   172.16.99.17   ip-192-168-100-171.us-east-2.compute.internal
mydeployment-665d88877-n5qbj  1/1    Running   0          3s    172.16.71.70   ip-192-168-100-118.us-east-2.compute.internal
```

From the above we can conclude that pod's are getting scheduled in the node where we have label set as "env=dev"

3. Resource Management:

```
[ec2-user@master kubernetes-installation]$ cat redis.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: redis
    name: redis
spec:
  replicas: 1
  selector:
    matchLabels:
      app: redis
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: redis
    spec:
      containers:
        - image: redis:3.2-alpine
          name: redis
          resources:
            limits:
              cpu: 1
              memory: 200Mi
            requests:
              cpu: 80m
              memory: 20Mi
[ec2-user@master kubernetes-installation]$ kubectl apply -f redis.yaml
deployment.apps/redis configured
[ec2-user@master kubernetes-installation]$
```

```
[ec2-user@master kubernetes-installation]$ kubectl get pods -l app=redis -o wide
NAME           READY   STATUS    RESTARTS   AGE   IP      NODE   NOMINATED NODE   READINESS GATES
redis-749c46d5bd-5nfqf  1/1    Running   0          79s   10.36.0.1   worker2   <none>        <none>
[ec2-user@master kubernetes-installation]$ kubectl get pods -l app=redis -o yaml
```

```
apiVersion: v1
items:
- apiVersion: v1
  kind: Pod
  metadata:
    creationTimestamp: "2020-08-12T02:52:46Z"
```

```
uid: 5/de/091-6/e6-4590-8a10-b6491a1239d2
resourceVersion: "51103"
selfLink: /api/v1/namespaces/default/pods/redis-749c46d5bd-5nfqf
uid: ed6cd0fa-881a-40a3-9466-8aac3d4b39b6
spec:
  containers:
    - image: redis:3.2-alpine
      imagePullPolicy: IfNotPresent
      name: redis
  resources:
    limits:
      cpu: "1"
      memory: 200Mi
    requests:
      cpu: 80m
      memory: 20Mi
  terminationMessagePath: /dev/termination-log
  terminationMessagePolicy: File
  volumeMounts:
    - mountPath: /var/run/secrets/kubernetes.io/serviceaccount
      name: default-token-c9tfm
      readOnly: true
```

4. DaemonSets and Manually Scheduled Pods :

```
[ec2-user@master kubernetes-installation]$ cat deamon.yaml
apiVersion: apps/v1
kind: "DaemonSet"
metadata:
  labels:
    app: nginx
    ssd: "true"
  name: nginx-fast-storage
spec:
  selector:
    matchLabels:
      app: nginx
      ssd: "true"
  template:
    metadata:
      labels:
        app: nginx
        ssd: "true"
  spec:
    nodeSelector:
      ssd: "true"
    containers:
      - name: nginx
        image: nginx:1.10.0
[ec2-user@master kubernetes-installation]$ kubectl create -f deamon.yaml
daemonset.apps/nginx-fast-storage created
```

```
[ec2-user@master kubernetes-installation]$ kubectl get ds
NAME           DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
nginx-fast-storage   0         0         0         0            0           ssd=true       9s
[ec2-user@master kubernetes-installation]$ kubectl get nodes -l ssd=true
No resources found in default namespace.
[ec2-user@master kubernetes-installation]$ kubectl label node worker1 ssd=true
node/worker1 labeled
[ec2-user@master kubernetes-installation]$ kubectl get nodes -l ssd=true
NAME   STATUS   ROLES   AGE   VERSION
worker1   Ready   <none>   2d   v1.18.6
[ec2-user@master kubernetes-installation]$ kubectl get ds
NAME           DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
nginx-fast-storage   1         1         1         1            1           ssd=true       64s
[ec2-user@master kubernetes-installation]$ kubectl get pods -l ssd=true
NAME           READY   STATUS   RESTARTS   AGE
nginx-fast-storage-d7gkq   1/1     Running   0          43s
[ec2-user@master kubernetes-installation]$ kubectl get pods -l ssd=true -o wide
NAME           READY   STATUS   RESTARTS   AGE   IP           NOMINATED NODE   READINESS GATES
nginx-fast-storage-d7gkq   1/1     Running   0          50s   10.44.0.4   worker1   <none>
[ec2-user@master kubernetes-installation]$ kubectl get nodes -l ssd=true
NAME   STATUS   ROLES   AGE   VERSION
worker1   Ready   <none>   2d   v1.18.6
```

```
[ec2-user@master kubernetes-installation]$ kubectl label node worker2 ssd=true
node/worker2 labeled
[ec2-user@master kubernetes-installation]$ kubectl get pods -l ssd=true -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP          NODE      NOMINATED NODE   READINESS
GATES
nginxx-fast-storage-d7gkq   1/1     Running   0          110s   10.44.0.4   worker1   <none>        <none>
nginxx-fast-storage-w98tw   0/1     ContainerCreating   0          8s     <none>       worker2   <none>        <none>
[ec2-user@master kubernetes-installation]$ kubectl get ds
NAME            DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
nginxx-fast-storage   2         2         2         2           2           ssd=true   2m52s
[ec2-user@master kubernetes-installation]$ kubectl label node worker2 ssd-
```

```
[ec2-user@master kubernetes-installation]$ kubectl label node worker2 ssd-
node/worker2 labeled
[ec2-user@master kubernetes-installation]$ kubectl get pods -l ssd=true -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP          NODE      NOMINATED NODE   READINESS
GATES
nginxx-fast-storage-d7gkq   1/1     Running   0          2m42s  10.44.0.4   worker1   <none>        <none>
nginxx-fast-storage-w98tw   0/1     Terminating   0          60s    10.36.0.2   worker2   <none>        <none>
[ec2-user@master kubernetes-installation]$
```

5. Displaying Scheduler Events:

This section helps us more in checking the events happening at the scheduler level.

Listing the scheduler pod.

```
[root@master1 ~]# kubectl get pods -n kube-system|grep -i schedul
kube-scheduler-master1   1/1     Running   1           19h
```

Checking the logs

```
[root@master1 ~]# kubectl logs kube-scheduler-master1 -f -n kube-system
```

To get all the events in the project
kubectl get events -n <project name>

To Delete all the pods in the project/namespace

```
kubectl delete pods --all -n mynginx
```

Watch events as they are appearing in real time:

```
kubectl get events -w
```

The location of a systemd service scheduler pod:

```
/var/log/kube-scheduler.log
```

7. Application lifecycle management

Installing and configuring docker in the workstation:

```
yum install docker -y
systemctl enable docker
systemctl restart docker
```

```
[root@workstation ~]# systemctl status docker
● docker.service - Docker Application Container Engine
  Loaded: loaded (/usr/lib/systemd/system/docker.service; enabled; vendor preset: disabled)
  Active: active (running) since Sat 2023-12-16 02:29:10 UTC; 20min ago
    Docs: https://docs.docker.com
 Process: 3800 ExecStartPre=/usr/libexec/docker/docker-setup-runtimes.sh (code=exited, status=0/SUCCESS)
S) Process: 3799 ExecStartPre=/bin/mkdir -p /run/docker (code=exited, status=0/SUCCESS)
Main PID: 3803 (dockerd)
  Tasks: 7
 Memory: 20.3M
 CGroup: /system.slice/docker.service
         └─3803 /usr/bin/dockerd -H fd:// --containerd=/run/containerd/containerd.sock --default-u...
Dec 16 02:29:09 workstation dockerd[3803]: time="2023-12-16T02:29:09.8994841Z" level=info msg="...rpc
Dec 16 02:29:09 workstation dockerd[3803]: time="2023-12-16T02:29:09.948628977Z" level=warning ms...ht"
Dec 16 02:29:09 workstation dockerd[3803]: time="2023-12-16T02:29:09.949138521Z" level=warning ms...ce"
Dec 16 02:29:09 workstation dockerd[3803]: time="2023-12-16T02:29:09.949654252Z" level=info msg="...t."
Dec 16 02:29:10 workstation dockerd[3803]: time="2023-12-16T02:29:10.186262647Z" level=info msg="...ss"
Dec 16 02:29:10 workstation dockerd[3803]: time="2023-12-16T02:29:10.232813777Z" level=info msg="...e."
Dec 16 02:29:10 workstation dockerd[3803]: time="2023-12-16T02:29:10.249123824Z" level=info msg="...25
Dec 16 02:29:10 workstation dockerd[3803]: time="2023-12-16T02:29:10.249612112Z" level=info msg="...on"
Dec 16 02:29:10 workstation systemd[1]: Started Docker Application Container Engine.
Dec 16 02:29:10 workstation dockerd[3803]: time="2023-12-16T02:29:10.276199375Z" level=info msg="...ck"
Hint: Some lines were ellipsized, use -l to show in full.
[root@workstation ~]#
```

Creating a custom image & pushing it to Dockerhub

```
[ec2-user@master ~]$ cat Dockerfile
FROM httpd
RUN echo "Hi I am V1" > /usr/local/apache2/htdocs/index.html
[ec2-user@master ~]$ sudo docker build -t "shan5a6/myimage:v1" .
Sending build context to Docker daemon 3.639 MB
Step 1/2 : FROM httpd
--> a6ea92c35c43
Step 2/2 : RUN echo "Hi I am V1" > /usr/local/apache2/htdocs/index.html
--> Using cache
--> baaf3cf68032
Successfully built baaf3cf68032
[ec2-user@master ~]$ sudo docker login
Login with your Docker ID to push and pull images from Docker Hub. If you don't have a Docker ID, head over to https://hub.docker.com to create one.
Username: (shan5a6): shan5a6
Password:
Login Succeeded
[ec2-user@master ~]$ docker push shan5a6/myimage:v1
Got permission denied while trying to connect to the Docker daemon socket at unix:///var/run/docker.sock: Post http://%2Fv
ar%2Frun%2Fdocker.sock/v1.26/images/shan5a6/myimage/push?tag=v1: dial unix /var/run/docker.sock: connect: permission denied
[ec2-user@master ~]$ sudo docker push shan5a6/myimage:v1
The push refers to a repository [docker.io/shan5a6/myimage]
7cc39fb2281c: Pushed
88b680b1fdfc: Mounted from library/httpd
843c3701e622: Mounted from library/httpd
3ba8a4f66da2: Mounted from library/httpd
c865989f86f7: Mounted from library/httpd
d0f104dc0a1f: Mounted from library/httpd
v1: digest: sha256:59fe38c4d1180e7ffb7c5ebec68e2e702ae8243e568fc428a049ffefaebe377 size: 1574
[ec2-user@master ~]$
```

Creating deployment with custom image:

Note: Make sure you are using "kubectl create -f deploy.yaml --record" for creating the deployment

```
[ec2-user@master ~]$ cat deploy.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mydeployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: mydeployment
  template:
    metadata:
      name: mydeployment
      labels:
        app: mydeployment
    spec:
      containers:
        - image: shan5a6/myimage:v1
          name: app
[ec2-user@master ~]$ kubectl apply -f deploy.yaml
deployment.apps/mydeployment created
[ec2-user@master ~]$ kubectl get deploy
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
mydeployment   0/3     3            0           5s
[ec2-user@master ~]$ kubectl get pods
NAME             READY   STATUS    RESTARTS   AGE
mydeployment-598689ff9f-bl6m6   1/1     Running   0          11s
mydeployment-598689ff9f-h2gz9   1/1     Running   0          11s
mydeployment-598689ff9f-q8m29   1/1     Running   0          11s
```

Creating service and verifying the current application status:

```
[ec2-user@master ~]$ kubectl create -f deploy.yaml --record
deployment.apps/mydeployment created
[ec2-user@master ~]$ kubectl get pods
NAME             READY   STATUS    RESTARTS   AGE
mydeployment-598689ff9f-kjdkj   1/1     Running   0          49s
mydeployment-598689ff9f-v9xb4   1/1     Running   0          49s
mydeployment-598689ff9f-z5wp6   1/1     Running   0          49s
nginx-fast-storage-x8r98       1/1     Running   0          10m
[ec2-user@master ~]$ kubectl get -o yaml deploy mydeployment | grep -i image
  f:image: {}
  f:imagePullPolicy: {}
  - image: shan5a6/myimage:v1
    imagePullPolicy: IfNotPresent
[ec2-user@master ~]$ kubectl expose deploy mydeployment --port 80 --target-port 80 --name mydeploysvc
service/mydeploysvc exposed
[ec2-user@master ~]$ kubectl describe svc mydeploysvc
Name:           mydeploysvc
Namespace:      default
Labels:         <none>
Annotations:   <none>
Selector:       app=mydeployment
Type:          ClusterIP
IP:            10.110.38.173
Port:          <unset>  80/TCP
TargetPort:    80/TCP
Endpoints:     10.36.0.1:80,10.36.0.2:80,10.44.0.3:80
Session Affinity:  None
Events:        <none>
[ec2-user@master ~]$
```

```
[ec2-user@master ~]$ curl http://10.110.38.173  
Hi I am V1  
[ec2-user@master ~]$ while true  
> do  
> curl http://10.110.38.173  
> sleep 5  
> done  
  
Hi I am V1  
Hi I am V1
```

↓
SUC IP

Note: If you are getting loadbalncer url then use the same for hitting and displaying application status. In this case

Creating customized image for the new release:

```
[ec2-user@master ~]$ cat Dockerfile
FROM httpd
RUN echo "Hi I am V2" > /usr/local/apache2/htdocs/index.html
[ec2-user@master ~]$ sudo docker build -t "shan5a6/myimage:v2" .
Sending build context to Docker daemon 3.64 MB
Step 1/2 : FROM httpd
--> a6ea92c35c43
Step 2/2 : RUN echo "Hi I am V2" > /usr/local/apache2/htdocs/index.html
--> Running in 12bd4bfal388

--> 0ale938f1efe
Removing intermediate container 12bd4bfal388
Successfully built 0ale938f1efe
[ec2-user@master ~]$ sudo docker push shan5a6/myimage:v2
The push refers to a repository [docker.io/shan5a6/myimage]
317fa498a6a9: Pushed
38b680b1fd9c: Layer already exists
343c3701e622: Layer already exists
3ba8a4f66ba2: Layer already exists
c865989f86f7: Layer already exists
d0f104dc0a1f: Layer already exists
v2: digest: sha256:e959e13896dde4781b9613fb3ebfce94176365231264c475459eec6031b0a9cc
[ec2-user@master ~]$
```

```
[ec2-user@master ~]$ kubectl scale deploy mydeployment --replicas=6
deployment.apps/mydeployment scaled
[ec2-user@master ~]$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
mydeployment-598689ff9f-b7dhq  1/1     Running   0          51s
mydeployment-598689ff9f-fxj2r  1/1     Running   0          2s
mydeployment-598689ff9f-jqfgb  1/1     Running   0          53s
mydeployment-598689ff9f-jzdp9  1/1     Running   0          2s
mydeployment-598689ff9f-kdv45  1/1     Running   0          50s
mydeployment-598689ff9f-nttks  1/1     Running   0          2s
nginx-fast-storage-x8r98      1/1     Running   0          22m
[ec2-user@master ~]$ vi deploy.yaml
[ec2-user@master ~]$ cat deploy.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mydeployment
spec:
  replicas: 6
  selector:
    matchLabels:
      app: mydeployment
  template:
    metadata:
      name: mydeployment
      labels:
        app: mydeployment
    spec:
      containers:
        - image: shan5a6/myimage:v2
          name: app
[ec2-user@master ~]$
```

Take two terminals and monitor the deployment changes like below:

```

Hi I am V1 [ec2-user@master ~]$ kubectl apply -f deploy.yaml --record
Hi I am V1 deployment.apps/mydeployment configured
Hi I am V1 [ec2-user@master ~]$ kubectl get pods -w
Hi I am V1 NAME READY STATUS RESTARTS AGE
Hi I am V1 mydeployment-598689ff9f-b7dhq 0/1 Terminating 0 2m47s
Hi I am V1 mydeployment-598689ff9f-fxj2r 0/1 Terminating 0 118s
Hi I am V1 mydeployment-598689ff9f-jqfgb 1/1 Terminating 0 2m49s
Hi I am V1 mydeployment-598689ff9f-jzdp9 0/1 Terminating 0 118s
Hi I am V1 mydeployment-598689ff9f-kdv45 1/1 Terminating 0 2m46s
Hi I am V1 mydeployment-598689ff9f-nttks 0/1 Terminating 0 118s
Hi I am V1 mydeployment-6ddd66b4c4-fg72m 1/1 Running 0 5s
Hi I am V1 mydeployment-6ddd66b4c4-hbskn 1/1 Running 0 2s
Hi I am V1 mydeployment-6ddd66b4c4-nngdr 1/1 Running 0 3s
Hi I am V1 mydeployment-6ddd66b4c4-rbsv7 1/1 Running 0 4s
Hi I am V1 mydeployment-6ddd66b4c4-v2c7s 1/1 Running 0 5s
Hi I am V1 mydeployment-6ddd66b4c4-v9pr9 1/1 Running 0 24m
Hi I am V1 nginx-fast-storage-x8r98 1/1 Running 0
Hi I am V2 mydeployment-598689ff9f-kdv45 0/1 Terminating 0 2m47s
Hi I am V2 mydeployment-598689ff9f-jqfgb 0/1 Terminating 0 2m50s
Hi I am V2 mydeployment-598689ff9f-nttks 0/1 Terminating 0 2m5s
Hi I am V2 mydeployment-598689ff9f-nttks 0/1 Terminating 0 2m5s
Hi I am V2 mydeployment-598689ff9f-jqfgb 0/1 Terminating 0 2m56s
Hi I am V2 mydeployment-598689ff9f-jqfgb 0/1 Terminating 0 2m56s
Hi I am V2 mydeployment-598689ff9f-fxj2r 0/1 Terminating 0 2m5s
Hi I am V2 mydeployment-598689ff9f-fxj2r 0/1 Terminating 0 2m5s
Hi I am V2 mydeployment-598689ff9f-jzdp9 0/1 Terminating 0 2m5s
Hi I am V2 mydeployment-598689ff9f-jzdp9 0/1 Terminating 0 2m5s
Hi I am V2 mydeployment-598689ff9f-kdv45 0/1 Terminating 0 2m53s
Hi I am V2 mydeployment-598689ff9f-kdv45 0/1 Terminating 0 2m53s
Hi I am V2 mydeployment-598689ff9f-b7dhq 0/1 Terminating 0 2m54s
Hi I am V2 mydeployment-598689ff9f-b7dhq 0/1 Terminating 0 2m54s
^C [ec2-user@master ~]$ kubectl get pods -w

```

Rollback to the previous version:

```

[ec2-user@master ~]$ while true; do curl http://10.110.38.173; sleep 5; done
[ec2-user@master ~]$ kubectl rollout history deployment mydeployment
error: required resource not specified
[ec2-user@master ~]$ kubectl rollout history deployment mydeployment
REVISION CHANGE-CAUSE
3   kubectl apply --filename=deploy.yaml --record=true
4   kubectl apply --filename=deploy.yaml --record=true
[ec2-user@master ~]$ kubectl rollout undo deployment mydeployment --to-revision=3
deployment.apps/mydeployment rolled back
[ec2-user@master ~]$ kubectl get pods
NAME READY STATUS RESTARTS AGE
Hi I am V1 mydeployment-598689ff9f-69dgm 1/1 Running 0 6s
Hi I am V1 mydeployment-598689ff9f-bx9z8 1/1 Running 0 8s
Hi I am V1 mydeployment-598689ff9f-dqrjl 1/1 Running 0 6s
Hi I am V1 mydeployment-598689ff9f-lfgsv 1/1 Running 0 8s
Hi I am V1 mydeployment-598689ff9f-sfqrb 1/1 Running 0 7s
Hi I am V1 mydeployment-598689ff9f-xmq2r 1/1 Running 0 8s
Hi I am V1 mydeployment-6ddd66b4c4-fg72m 0/1 Terminating 0 7m56s
Hi I am V1 mydeployment-6ddd66b4c4-v2c7s 0/1 Terminating 0 7m56s
Hi I am V1 nginx-fast-storage-x8r98 1/1 Running 0 32m
[ec2-user@master ~]$ kubectl get pods
NAME READY STATUS RESTARTS AGE
mydeployment-598689ff9f-69dgm 1/1 Running 0 40s
mydeployment-598689ff9f-bx9z8 1/1 Running 0 42s
mydeployment-598689ff9f-dqrjl 1/1 Running 0 40s
mydeployment-598689ff9f-lfgsv 1/1 Running 0 42s
mydeployment-598689ff9f-sfqrb 1/1 Running 0 41s

```

8. Configuring an Application for High Availability and Scale

1. Configmaps :

Creating configmap:

```
[ec2-user@master ~]$ cat configmap.yaml
apiVersion: v1
kind: ConfigMap
metadata:
  name: appconfig
data:
  key1: value1
  key2: value2
[ec2-user@master ~]$ kubectl create -f configmap.yaml
configmap/appconfig created
[ec2-user@master ~]$ kubectl describe configmap appconfig
Name:           appconfig
Namespace:      default
Labels:         <none>
Annotations:   <none>

Data
=====
key1:
-----
value1
key2:
-----
value2
Events:  <none>
[ec2-user@master ~]$
```

Creating pod with configmap

```
[ec2-user@master ~]$ cat configmap-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: configmap-pod
spec:
  containers:
    - name: app-container
      image: busybox:1.28
      command: ['sh', '-c', "echo ${MY_VAR} && sleep 3600"]
    env:
      - name: MY_VAR
        valueFrom:
          configMapKeyRef:
            name: appconfig
            key: key1
[ec2-user@master ~]$ kubectl apply -f configmap-pod.yaml
pod/configmap-pod created
[ec2-user@master ~]$ kubectl get pods | grep -i app-container
[ec2-user@master ~]$ kubectl get pods
NAME             READY   STATUS    RESTARTS   AGE
configmap-pod   1/1     Running   0          27s
```

Verifying details inside the pod:

```
[ec2-user@master ~]$ kubectl exec -it configmap-pod sh  
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version.  
MAND instead.  
/ # env|grep -i MYVAR  
/ # env|grep -i MY  
MY_VAR=value1  
MYDEPLOY SVC SERVICE_HOST=10.110.38.173  
MYDEPLOY SVC SERVICE_PORT=80  
MYDEPLOY SVC PORT=tcp://10.110.38.173:80  
MYDEPLOY SVC PORT_80_TCP_ADDR=10.110.38.173  
MYDEPLOY SVC PORT_80_TCP_PORT=80  
MYDEPLOY SVC PORT_80_TCP_PROTO=tcp  
MYDEPLOY SVC PORT_80_TCP=tcp://10.110.38.173:80  
[ec2-user@master ~]$ kubectl describe configmap appconfig  
Name:           appconfig  
Namespace:      default  
Labels:         <none>
```

2. Creating a secret

Let's create the secret

```
[ec2-user@master ~]$ cat secret.yaml  
apiVersion: v1  
kind: Secret  
metadata:  
  name: appsecret  
stringData:  
  subject: "kubernetes"  
  batchno: "3"  
[ec2-user@master ~]$ kubectl create -f secret.yaml  
secret/appsecret created  
[ec2-user@master ~]$ kubectl describe secret appsecret  
Name:           appsecret  
Namespace:      default  
Labels:         <none>  
Annotations:    <none>  
  
Type:  Opaque  
  
Data  
====  
batchno: 1 bytes  
subject: 10 bytes  
[ec2-user@master ~]$ kubectl get -o yaml secret appsecret  
apiVersion: v1  
data:  
  batchno: Mw==  
  subject: a3ViZXJuZXRlcw==  
kind: Secret  
metadata:  
  creationTimestamp: "2020-08-13T02:58:09Z"  
  managedFields:  
  - apiVersion: v1  
    fieldsType: FieldsV1
```

Encrypted data

Creating a pod:

```
[ec2-user@master ~]$ cat secret-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: secret-pod
spec:
  containers:
  - name: app-container
    image: busybox
    command: ['sh', '-c', "echo Hello, Kubernetes! && sleep 3600"]
    env:
      - name: MY_SUB
        valueFrom:
          secretKeyRef:
            name: appsecret
            key: subject
[ec2-user@master ~]$ kubectl create -f secret-pod.yaml
pod/secret-pod created
[ec2-user@master ~]$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
configmap-pod  1/1     Running   0          10m
mydeployment-598689ff9f-69dgm  1/1     Running   0          32m
mydeployment-598689ff9f-bx9z8  1/1     Running   0          32m
mydeployment-598689ff9f-dqrjl  1/1     Running   0          32m
mydeployment-598689ff9f-lfgsv  1/1     Running   0          32m
mydeployment-598689ff9f-sfqrb  1/1     Running   0          32m
mydeployment-598689ff9f-xmq2r  1/1     Running   0          32m
nginx-fast-storage-x8r98       1/1     Running   0          65m
secret-pod      1/1     Running   0          4s
[ec2-user@master ~]$
```

```
[ec2-user@master ~]$ kubectl exec -it secret-pod sh
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future
MAND instead.
/ # env|grep -i my
MY_SUB=kubernetes
MYDEPLOYSVCSERVICE_HOST=10.110.38.173
MYDEPLOYSVCSERVICE_PORT=tcp://10.110.38.173:80
```

Encrypting & decrypting the secrets:

```
[ec2-user@master ~]$ echo "dvsbatch3" |base64
ZHZZYmF0Y2gzCg==
[ec2-user@master ~]$ echo "ZHZZYmF0Y2gzCg==" |base64 -d
dvsbatch3
[ec2-user@master ~]$
```

9. Storage

In Eks we can make use of aws ebs volume by default in the cluster we will not have the storage class we have to enable it. Please do execute the below to enable ebs storage class in eks cluster.

Storage Management::

<https://www.stacksimplify.com/aws-eks/kubernetes-storage/install-aws-ebs-csi-driver-on-aws-eks-for-persistent-storage/>

Ref: Git Link

<https://github.com/kubernetes-sigs/aws-ebs-csi-driver.git>

```
POLICY_ARN=$(aws iam list-policies --query 'Policies[?PolicyName==`AmazonEBSCSIDriverPolicy`].Arn' --output text)
NODE_ROLE_NAME=$(kubectl get configmap aws-auth -n kube-system -o yaml |grep rolearn|awk '{print $2}'|awk -F'/' '{print $2}')
aws iam attach-role-policy --policy-arn $POLICY_ARN --role-name $NODE_ROLE_NAME
```

```
[root@workstation myobjects]# POLICY_ARN=$(aws iam list-policies --query 'Policies[?PolicyName==`AmazonEBSCSIDriverPolicy`].Arn' --output text)
NODE_ROLE_NAME=$(kubectl get configmap aws-auth -n kube-system -o yaml |grep rolearn|awk '{print $2}'|awk -F'/' '{print $2}')
aws iam attach-role-policy --policy-arn $POLICY_ARN --role-name $NODE_ROLE_NAME
[root@workstation myobjects]# NODE_ROLE_NAME=$(kubectl get configmap aws-auth -n kube-system -o yaml |grep rolearn|awk '{print $2}'|awk -F'/' '{print $2}')
[root@workstation myobjects]# aws iam attach-role-policy --policy-arn $POLICY_ARN --role-name $NODE_ROLE_NAME
[root@workstation myobjects]#
[root@workstation myobjects]# echo $POLICY_ARN
arn:aws:iam::aws:policy/service-role/AmazonEBSCSIDriverPolicy
[root@workstation myobjects]# echo $NODE_ROLE_NAME
eksctl-dvsekscluster-nodegroup-nod-NodeInstanceRole-PByQJHwQfDyk
[root@workstation myobjects]#
```

Above commands attach the AmazonEBSCSIDriverPolicy to the node role "eksctl-dvsekscluster-nodegroup-nod-NodeInstanceRole-PByQJHwQfDyk"

Install EBS CSI Driver

```
kubectl apply -k "github.com/kubernetes-sigs/aws-ebs-csi-driver/deploy/kubernetes/overlays/stable/?ref=master"
```

```
[root@workstation myobjects]# kubectl apply -k "github.com/kubernetes-sigs/aws-ebs-csi-driver/deploy/kubernetes/overlays/stable/?ref=master"
serviceaccount/ebs-csi-controller-sa created
serviceaccount/ebs-csi-node-sa created
role.rbac.authorization.k8s.io/ebs-csi-leases-role created
clusterrole.rbac.authorization.k8s.io/ebs-csi-node-role created
clusterrole.rbac.authorization.k8s.io/ebs-external-attacher-role created
clusterrole.rbac.authorization.k8s.io/ebs-external-provisioner-role created
clusterrole.rbac.authorization.k8s.io/ebs-external-resizer-role created
clusterrole.rbac.authorization.k8s.io/ebs-external-snapshotter-role created
rolebinding.rbac.authorization.k8s.io/ebs-csi-leases-rolebinding created
clusterrolebinding.rbac.authorization.k8s.io/ebs-csi-attacher-binding created
clusterrolebinding.rbac.authorization.k8s.io/ebs-csi-node-getter-binding created
clusterrolebinding.rbac.authorization.k8s.io/ebs-csi-provisioner-binding created
```

```
kubectl create -f - <<EOF
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: ebs-sc
provisioner: ebs.csi.aws.com
volumeBindingMode: WaitForFirstConsumer
EOF
```

```
[root@workstation myobjects]# kubectl create -f - <<EOF
> apiVersion: storage.k8s.io/v1
> kind: StorageClass
> metadata:
>   name: ebs-sc
>   provisioner: ebs.csi.aws.com
>   volumeBindingMode: WaitForFirstConsumer
> EOF
storageclass.storage.k8s.io/ebs-sc created
[root@workstation myobjects]# kubectl get sc
NAME      PROVISIONER      RECLAIMPOLICY      VOLUMEBINDINGMODE      ALLOWVOLUMEEXPANSION
ebs-sc    ebs.csi.aws.com  Delete            WaitForFirstConsumer  false
gp2 (default)  kubernetes.io/aws-ebs  Delete            WaitForFirstConsumer  false
4h
[root@workstation myobjects]#
```

Now from the above we conclude that storage class ebs-sc is created let's create the PVC and then the pod then verify if we are getting the pod with storage or not.

creating a persistent-volume claim:

```
kubectl create -f - <<EOF
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: ebs-claim
spec:
  accessModes:
    - ReadWriteOnce
  storageClassName: gp2
  resources:
```

DVS Technologies, Opp Home Town, Beside Biryani Zone, Marathahalli, Bangalore Phone: 9632558585 Mobile: 8892499499 Mail : dvs.training@gmail.com Web: www.dvstechnologies.in

```
requests:  
  storage: 4Gi  
EOF
```

```
*# [root@workstation myobjects]# kubectl create -f - <<EOF  
> apiVersion: v1  
> kind: PersistentVolumeClaim  
> metadata:  
>   name: ebs-claim  
> spec:  
>   accessModes:  
>     - ReadWriteOnce  
>   storageClassName: gp2  
>   resources:  
>     requests:  
>       storage: 4Gi  
> EOF  
  
persistentvolumeclaim/ebs-claim created  
[root@workstation myobjects]# [root@workstation myobjects]# kubectl get pvc  
NAME      STATUS    VOLUME   CAPACITY   ACCESS MODES  STORAGECLASS   AGE  
ebs-claim  Pending          gp2           7s  
[root@workstation myobjects]# kubectl get pv  
No resources found  
[root@workstation myobjects]#
```

Creating the consumer to consume the pod:
pod.yaml

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: app  
spec:  
  containers:  
    - name: app  
      image: centos  
      command: ["/bin/sh"]  
      args: ["-c", "while true; do echo $(date -u) >> /data/out.txt; sleep 5; done"]  
      volumeMounts:  
        - name: persistent-storage  
          mountPath: /data  
  volumes:  
    - name: persistent-storage  
      persistentVolumeClaim:  
        claimName: ebs-claim
```

```
[root@workstation myobjects]# cat pod-with-pv.yaml
apiVersion: v1
kind: Pod
metadata:
  name: app
spec:
  containers:
  - name: app
    image: centos
    command: ["/bin/sh"]
    args: ["-c", "while true; do echo $(date -u) >> /data/out.txt; sleep 5; done"]
    volumeMounts:
    - name: persistent-storage
      mountPath: /data
  volumes:
  - name: persistent-storage
    persistentVolumeClaim:
      claimName: ebs-claim

[root@workstation myobjects]# kubectl apply -f pod-with-pv.yaml
pod/app created
[root@workstation myobjects]# kubectl get pod
NAME     READY   STATUS    RESTARTS   AGE
app     1/1     Running   0          26s
[root@workstation myobjects]# kubectl get pv
```

Once we got our pod up & running we can see that we have our pv & pvc got created & bounded

```
Pod/app created
[root@workstation myobjects]# kubectl get pod
NAME     READY   STATUS    RESTARTS   AGE
app     1/1     Running   0          26s
[root@workstation myobjects]# kubectl get pv
NAME           CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM
STORAGECLASS   REASON    AGE
pvc-02fb9cb7-2ecd-47a5-b65c-e930elc32d19  4Gi        RWO          Delete   Bound   default/
ebs-claim      gp2       28s
[root@workstation myobjects]# kubectl get pvc
NAME           STATUS   VOLUME                                     CAPACITY   ACCESS MODES   STORAGECLASS
AGE
ebs-claim      Bound    pvc-02fb9cb7-2ecd-47a5-b65c-e930elc32d19  4Gi        RWO          gp2
105s
[root@workstation myobjects]#
```

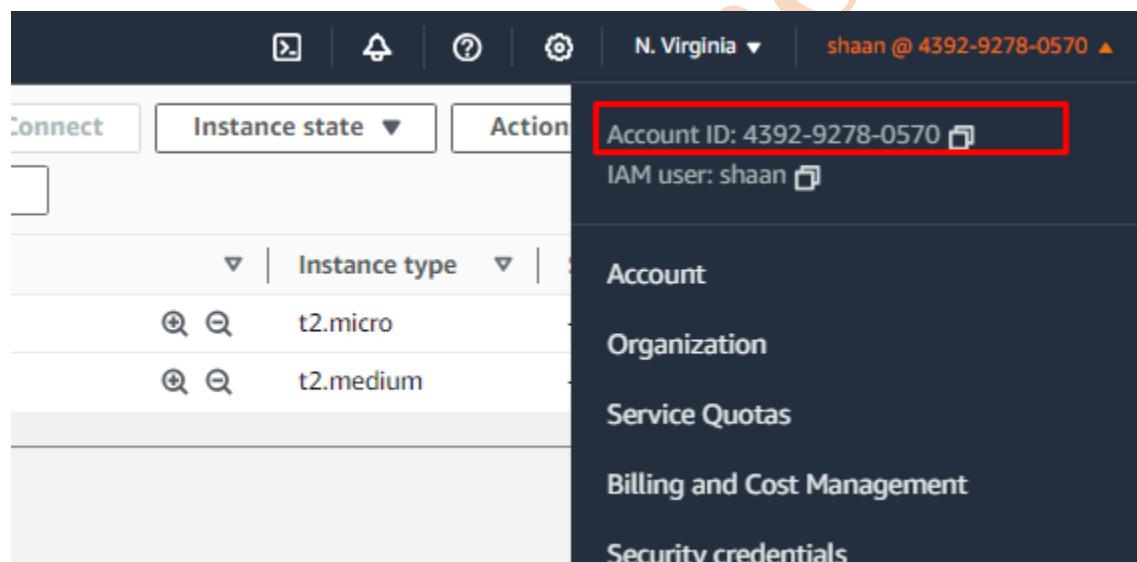
Name	Volume ID	Type	Size	IOPS	Throughput	Snapshot	Created
dvsekscluster-...	vol-04f6c363fedde150e	gp3	80 GiB	3000	125	snap-01a44b8...	2023/12/24 05:30 GMT+4
dvsekscluster-...	vol-05685ebfee0aee9ca	gp3	80 GiB	3000	125	snap-01a44b8...	2023/12/24 11:26 GMT+4
-	vol-0e52d96331a167734	gp2	4 GiB	100	-	-	2023/12/24 12:29 GMT+4

10 Cluster autoscalling:

Note: Make sure that you are changing the nodegroup ASG node type to atleast t2.medium otherwise node scheduler pod will not come up .

```
[root@workstation myobjects]# aws eks list-clusters
{
  "clusters": [
    "dvsekscluster"
  ]
}
[root@workstation myobjects]#
```

```
export CLUSTER_NAME="dvsekscluster" # GIVE YOUR CLUSTER NAME
export MIN=2
export MAX=4
export DESIRED=3
export ACCOUNT_ID=439292780570 # Give your account ID
```



```
[root@workstation myobjects]# export CLUSTER_NAME="dvsekscluster" # GIVE YOUR CLUSTER NAME
[root@workstation myobjects]# export MIN=2
[root@workstation myobjects]# export MAX=4
[root@workstation myobjects]# export DESIRED=3
[root@workstation myobjects]# export ACCOUNT_ID=439292780570 # Give your account ID
[root@workstation myobjects]#
```

```
aws autoscaling \
  describe-auto-scaling-groups \
```

```

--query "AutoScalingGroups[? Tags[? (Key=='eks:cluster-name') &&
Value=='$CLUSTER_NAME']].[AutoScalingGroupName, MinSize,
MaxSize,DesiredCapacity]" \
--output table
export ASG_NAME=$(aws autoscaling describe-auto-scaling-groups --query
"AutoScalingGroups[? Tags[? (Key=='eks:cluster-name') &&
Value=='$CLUSTER_NAME']].AutoScalingGroupName" --output text)

# increase max capacity up to 4
aws autoscaling \
    update-auto-scaling-group \
        --auto-scaling-group-name ${ASG_NAME} \
        --min-size $MIN \
        --desired-capacity $DESIRED \
        --max-size $MAX

# Check new values
aws autoscaling \
    describe-auto-scaling-groups \
        --query "AutoScalingGroups[? Tags[? (Key=='eks:cluster-name') &&
Value=='$CLUSTER_NAME']].[AutoScalingGroupName, MinSize,
MaxSize,DesiredCapacity]" \
        --output table

```

```

[root@workstation myobjects]# aws autoscaling \
>     describe-auto-scaling-groups \
>     --query "AutoScalingGroups[? Tags[? (Key=='eks:cluster-name') && Value=='$CLUSTER_NAME']].[AutoScalingGroupName, MinSize, MaxSize,DesiredCapacity]" \
>     --output table
-----
|           DescribeAutoScalingGroups           |
+---+---+---+---+
| eks-nodegroup1-7cc64ad4-5317-fbb0-f762-5a27e7fb2c6f | 2 | 2 | 2 |
+---+---+---+---+
[root@workstation myobjects]# export ASG_NAME=$(aws autoscaling describe-auto-scaling-groups --query "AutoScalingGroups[? Tags[? (Key=='eks:cluster-name') && Value=='$CLUSTER_NAME']].AutoScalingGroupName" --output text)
[root@workstation myobjects]# aws autoscaling \
>     update-auto-scaling-group \
>     --auto-scaling-group-name ${ASG_NAME} \
>     --min-size $MIN \
>     --desired-capacity $DESIRED \
>     --max-size $MAX
[root@workstation myobjects]# aws autoscaling \
>     describe-auto-scaling-groups \
>     --query "AutoScalingGroups[? Tags[? (Key=='eks:cluster-name') && Value=='$CLUSTER_NAME']].[AutoScalingGroupName, MinSize, MaxSize,DesiredCapacity]" \
>     --output table
-----
|           DescribeAutoScalingGroups           |
+---+---+---+---+
| eks-nodegroup1-7cc64ad4-5317-fbb0-f762-5a27e7fb2c6f | 2 | 4 | 3 |
+---+---+---+---+
[root@workstation myobjects]#

```

```

# Enabling IAM roles for service accounts on your cluster
eksctl utils associate-iam-oidc-provider \
--cluster $CLUSTER_NAME \
--approve

# Creating an IAM policy for your service account that will allow your CA pod to interact with
the autoscaling groups.

cat <<EoF >k8s-asg-policy.json
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Action": [
        "autoscaling:DescribeAutoScalingGroups",
        "autoscaling:DescribeAutoScalingInstances",
        "autoscaling:DescribeLaunchConfigurations",
        "autoscaling:DescribeTags",
        "autoscaling:SetDesiredCapacity",
        "autoscaling:TerminateInstanceInAutoScalingGroup",
        "ec2:DescribeLaunchTemplateVersions"
      ],
      "Resource": "*",
      "Effect": "Allow"
    }
  ]
}
EoF

```

```

+-----+-----+-----+
[root@workstation myobjects]# eksctl utils associate-iam-oidc-provider \
>   --cluster $CLUSTER_NAME \
>   --approve
2023-12-24 08:39:29 [!] will create IAM Open ID Connect provider for cluster "dvsekscluster" in "us-east-2"
2023-12-24 08:39:29 [v] created IAM Open ID Connect provider for cluster "dvsekscluster" in "us-east-2"
"
[root@workstation myobjects]# cat <<EoF >k8s-asg-policy.json
> {
>   "Version": "2012-10-17",
>   "Statement": [
>     {
>       "Action": [
>         "autoscaling:DescribeAutoScalingGroups",
>         "autoscaling:DescribeAutoScalingInstances",
>         "autoscaling:DescribeLaunchConfigurations",
>         "autoscaling:DescribeTags",
>         "autoscaling:SetDesiredCapacity",
>         "autoscaling:TerminateInstanceInAutoScalingGroup",
>         "ec2:DescribeLaunchTemplateVersions"
>       ],
>       "Resource": "*",
>       "Effect": "Allow"
>     }
>   ]
> }
> EoF
[root@workstation myobjects]#

```

```
aws iam create-policy --policy-name k8s-asg-policy --policy-document file:///\$\(pwd\)/k8s-asg-policy.json
```

```
eksctl create iamserviceaccount \
--name cluster-autoscaler \
--namespace kube-system \
--cluster ${CLUSTER_NAME} \
--attach-policy-arn "arn:aws:iam::${ACCOUNT_ID}:policy/k8s-asg-policy" \
--approve \
--override-existing-serviceaccounts
```

```
kubectl -n kube-system describe sa cluster-autoscaler
```

```
[root@workstation myobjects]# aws iam create-policy --policy-name k8s-asg-policy --policy-document file:///\$\(pwd\)/k8s-asg-policy.json
{
    "Policy": {
        "PolicyName": "k8s-asg-policy",
        "PolicyId": "ANPAWMR7D4ANJ5VF7ZKK",
        "Arn": "arn:aws:iam::439292780570:policy/k8s-asg-policy",
        "Path": "/",
        "DefaultVersionId": "v1",
        "AttachmentCount": 0,
        "PermissionsBoundaryUsageCount": 0,
        "IsAttachable": true,
        "CreateDate": "2023-12-24T08:41:25+00:00",
        "UpdateDate": "2023-12-24T08:41:25+00:00"
    }
}
[root@workstation myobjects]# eksctl create iamserviceaccount \
>   --name cluster-autoscaler \
>   --namespace kube-system \
>   --cluster ${CLUSTER_NAME} \
>   --attach-policy-arn "arn:aws:iam::${ACCOUNT_ID}:policy/k8s-asg-policy" \
>   --approve \
>   --override-existing-serviceaccounts
2023-12-24 08:41:33 [!] 1 iamserviceaccount (kube-system/cluster-autoscaler) was included (based on the include/exclude rules)
2023-12-24 08:41:33 [!] metadata of serviceaccounts that exist in Kubernetes will be updated, as --override-existing-serviceaccounts was set
2023-12-24 08:41:33 [!] 1 task:
  2 sequential sub-tasks:
    create IAM role for serviceaccount "kube-system/cluster-autoscaler",
    create serviceaccount "kube-system/cluster-autoscaler"
```

```
kubectl apply -f https://www.eksworkshop.com/beginner/080\_scaling/deploy\_ca.files/cluster-autoscaler-autodiscover.yaml
```

```
export AUTOSCALER_VERSION="1.21.2"
kubectl -n kube-system \
  set image deployment.apps/cluster-autoscaler \
  cluster-autoscaler=us.gcr.io/k8s-artifacts-prod/autoscaling/cluster-autoscaler:v${AUTOSCALER_VERSION}
```

```
[root@workstation myobjects]# kubectl apply -f https://www.eksworkshop.com/beginner/080_scaling/deploy_ca.files/cluster-autoscaler-autodiscover.yaml

export AUTOSCALER_VERSION="1.21.2"
kubectl -n kube-system \
    set image deployment.apps/cluster-autoscaler \
        cluster-autoscaler=us.gcr.io/k8s-artifacts-prod/autoscaling/cluster-autoscaler:v${AUTOSCALER_VERSION}
clusterrole.rbac.authorization.k8s.io/cluster-autoscaler created
role.rbac.authorization.k8s.io/cluster-autoscaler created
clusterrolebinding.rbac.authorization.k8s.io/cluster-autoscaler created
rolebinding.rbac.authorization.k8s.io/cluster-autoscaler created
deployment.apps/cluster-autoscaler created
[root@workstation myobjects]#
[root@workstation myobjects]# export AUTOSCALER_VERSION="1.21.2"
[root@workstation myobjects]# kubectl -n kube-system \
>     set image deployment.apps/cluster-autoscaler \
>         cluster-autoscaler=us.gcr.io/k8s-artifacts-prod/autoscaling/cluster-autoscaler:v${AUTOSCALER_VERSION}
deployment.apps/cluster-autoscaler image updated
[root@workstation myobjects]#
```

Make sure that you are changing the deployment & update the nodegroup vaules
kubectl edit deploy cluster-autoscaler -n kube-system

Edit the below section in the deployment

- --node-group-auto-discovery=asg:tag=k8s.io/cluster-autoscaler/enabled,k8s.io/cluster-autoscaler/dvsbatch3-eks-2{ update this one as per your cluster name }

Before:

```
deployment.apps/cluster-autoscaler image updated
[root@workstation myobjects]# kubectl edit deploy cluster-autoscaler -n kube-system
    prometheus.io/port: "8085"
    prometheus.io/scrape: "true"
    creationTimestamp: null
    labels:
        app: cluster-autoscaler
    spec:
        containers:
            - command:
                - ./cluster-autoscaler
                - --v=4
                - --stderrthreshold=info
                - --cloud-provider=aws
                - --skip-nodes-with-local-storage=false
                - --expander=least-waste
                - --node-group-auto-discovery=asg:tag=k8s.io/cluster-autoscaler/enabled,k8s.io/cluster-autoscaler/eksworkshop-eksct
                - --balance-similar-node-groups
                - --skip-nodes-with-system-pods=false
            image: us.gcr.io/k8s-artifacts-prod/autoscaling/cluster-autoscaler:v1.21.2
            imagePullPolicy: Always
            name: cluster-autoscaler
```

After:

```
[root@workstation myobjects]# kubectl edit deploy cluster-autoscaler -n kube-system
    prometheus.io/port: "8085"
    prometheus.io/scrape: "true"
  creationTimestamp: null
  labels:
    app: cluster-autoscaler
  spec:
    containers:
      command:
        - ./cluster-autoscaler
        - --v=4
        - --stderrthreshold=info
        - --cloud-provider=aws
        - --skip-nodes-with-local-storage=false
        - --expander=least-waste
        - --node-group-auto-discovery=asg:tag=k8s.io/cluster-autoscaler/enabled,k8s.io/cluster-autoscaler/dvsekscuster
        - --balance-similar-node-groups
        - --skip-nodes-with-system-pods=false
      image: us.gcr.io/k8s-artifacts-prod/autoscaling/cluster-autoscaler:v1.21.2
      imagePullPolicy: Always
      name: cluster-autoscaler
      resources:
        limits:
          cpu: 100m
          memory: 500Mi
        requests:
          cpu: 100m
```

Now we are done with our changes for creating cluster auto scaler

Let's test the setup.

```
[root@workstation myobjects]# kubectl get nodes
NAME           STATUS   ROLES   AGE     VERSION
ip-192-168-100-118.us-east-2.compute.internal Ready   <none>  83m   v1.27.7-eks-e71965b
ip-192-168-100-157.us-east-2.compute.internal Ready   <none>  12m   v1.27.7-eks-e71965b
ip-192-168-100-171.us-east-2.compute.internal Ready   <none>  7h20m  v1.27.7-eks-e71965b
[root@workstation myobjects]# vi deploy.yaml
[root@workstation myobjects]# cat deploy.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
  name: mydeployment
  namespace: myapplication
spec:
  replicas: 100
  selector:
    matchLabels:
      app: mydeployment
  strategy: {}
  template:
```

As of now we have 3 nodes and we are increasing our deployment values to 100 which means we are increasing load to our infrastructure as part of the cluster autoscaler it should scale the servers to max value in our case its 4

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
  name: mydeployment
  namespace: myapplication
```

```

spec:
  replicas: 1000
  selector:
    matchLabels:
      app: mydeployment
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
    labels:
      app: mydeployment
  spec:
    containers:
      - image: nginx
        name: nginx
        resources: {}
status: {}

```

```

[root@workstation myobjects]# cat deploy.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: mydeployment
    name: mydeployment
    namespace: myapplication
spec:
  replicas: 1000
  selector:
    matchLabels:
      app: mydeployment
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
    labels:
      app: mydeployment
  spec:
    containers:
      - image: nginx
        name: nginx
        resources: {}
status: {}
[root@workstation myobjects]# kubectl apply -f deploy.yaml
deployment.apps/mydeployment configured
[root@workstation myobjects]#

```

```

^[[B^C[root@workstation myobjects]# kubectl get nodes
NAME                               STATUS   ROLES   AGE     VERSION
ip-192-168-100-105.us-east-2.compute.internal Ready    <none>  11m    v1.27.7-eks-e71965b
ip-192-168-100-118.us-east-2.compute.internal Ready    <none>  112m   v1.27.7-eks-e71965b
ip-192-168-100-157.us-east-2.compute.internal Ready    <none>  40m    v1.27.7-eks-e71965b
ip-192-168-100-171.us-east-2.compute.internal Ready    <none>  7h49m   v1.27.7-eks-e71965b
[root@workstation myobjects]#

```

Hence we can say that cluster auto scaler is working fine.

11. Ingress

<https://kubernetes.github.io/ingress-nginx/deploy/#azure>

Installing Ingress: (Manually)

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.1.2/deploy/static/provider/cloud/deploy.yaml
```

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.1.3/deploy/static/provider/cloud/deploy.yaml
```

Make sure that you have below softwares in the workstation

Installing helm:

```
curl https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3 | bash
```

```
[root@workstation ~]# curl https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3 | bash
% Total    % Received % Xferd  Average Speed   Time     Time   Current
          Dload  Upload Total   Spent    Left Speed
100 11156  100 11156    0      0  63015      0 --:--:-- --:--:-- 63386
Downloaded https://get.helm.sh/helm-v3.8.1-linux-amd64.tar.gz
Verifying checksum... Done.
Preparing to install helm into /usr/local/bin
helm installed into /usr/local/bin/helm
[root@workstation ~]#
```

```
[root@workstation ~]# helm version
version.BuildInfo{Version:"v3.8.1", GitCommit:"5cb9af4b1b271d11d7a97a71df3ac337dd94ad37", GitTreeState:"clean", GoVersion:"go1.17.5"}
[root@workstation ~]#
```

Installing Ingress using helm:

<https://docs.microsoft.com/en-us/azure/aks/ingress-basic?tabs=azure-cli>

NAMESPACE=ingress-basic

```
helm repo add ingress-nginx https://kubernetes.github.io/ingress-nginx
helm repo update
```

```
helm install ingress-nginx ingress-nginx/ingress-nginx --create-namespace --namespace
$NAMESPACE
```

```
|root@workstation ~]# NAMESPACE=ingress-basic  
|root@workstation ~]# helm repo add ingress-nginx https://kubernetes.github.io/ingress-nginx  
"ingress-nginx" has been added to your repositories  
|root@workstation ~]# helm repo update  
Hang tight while we grab the latest from your chart repositories...  
...Successfully got an update from the "ingress-nginx" chart repository  
Update Complete. Happy Helming! I  
|root@workstation ~]#
```

```
[root@workstation ~]# helm install ingress-nginx ingress-nginx/ingress-nginx --create-namespace --namespace $NAMESPACE
NAME: ingress-nginx
LAST DEPLOYED: Mon Apr 11 06:47:42 2022
NAMESPACE: ingress-basic
STATUS: deployed
REVISION: 1
TEST SUITE: None
NOTES:
The ingress-nginx controller has been installed.
It may take a few minutes for the LoadBalancer IP to be available.
You can watch the status by running 'kubectl --namespace ingress-basic get services -o wide -w ingress-nginx-controller'

An example Ingress that makes use of the controller:
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example
  namespace: foo
spec:
  ingressClassName: nginx
```

```
[root@workstation myobjects]# kubectl get svc -n ingress-basic
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP
ingress-nginx-controller   LoadBalancer   10.100.149.246   ad7fbf9075cf64c9e90a252cda89978e-5
4684905.us-east-2.elb.amazonaws.com   80:3029/TCP, 443:31381/TCP   6n16m
ingress-nginx-controller-admission   ClusterIP    10.100.149.75   <none>
                                         443/TCP                  6h16m

[root@workstation myobjects]#
```

In the aws console:

The screenshot shows the AWS Management Console with the EC2 service selected. Under the 'Load Balancing' section, the 'Load balancers' tab is active. The main pane displays a table of load balancers with columns for Name, DNS name, State, VPC ID, and Availability. Two entries are visible: 'ad7fbf9075cf4c9e90a252cda89978e' (selected) and 'ad8c13e7436e04d36bec2dc1186f520d'. The details for the selected load balancer are shown in the bottom half of the screen, including its type (Classic), status (4 of 4 instances in service), VPC (vpc-0d160a3a9d5006215), creation date (December 24, 2023, 07:20 UTC+04:00), and scheme (Hosted zone).

Name	DNS name	State	VPC ID	Availability
<input checked="" type="checkbox"/> ad7fbf9075cf4c9e90a252cda89978e	ad7fbf9075cf4c9e90a25...	-	vpc-0d160a3a9d5006...	3 Av
<input type="checkbox"/> ad8c13e7436e04d36bec2dc1186f520d	ad8c13e7436e04d36bec2...	-	vpc-0d160a3a9d5006...	3 Av

Load balancer: ad7fbf9075cf4c9e90a252cda89978e

Details **Listeners** **Network mapping** **Security** **Health checks** **Target instances** **Monitoring** **Attributes** **Tags**

Details

Load balancer type	Status	VPC	Date created
Classic	4 of 4 instances in service	vpc-0d160a3a9d5006215	December 24, 2023, 07:20 (UTC+04:00)
Scheme	Hosted zone	Availability Zones	

Let's configure the application with ingress and test the routing:

DVS Technologies, Opp Home Town, Beside Biryani Zone, Marathahalli, Bangalore Phone: 9632558585 Mobile: 8892499499 Mail : dvs.training@gmail.com Web: www.dvstechnologies.in

Create two deployments:

Create two images & push to dockerhub:

App1:

Dockerfile

```
FROM httpd
RUN echo "<h1>Hi Team welcome to deployment1-app1</h1>" >
/usr/local/apache2/htdocs/index.html
```

```
docker build -t "shan5a6/deployment1-app1:v1.0".
```

```
[root@workstation myobjects]# vi Dockerfile
[root@workstation myobjects]# cat Dockerfile
FROM httpd
RUN echo "<h1>Hi Team welcome to deployment1-app1</h1>" > /usr/local/apache2/htdocs/index.html
[root@workstation myobjects]# docker build -t "shan5a6/deployment1-app1:v1.0".
Sending build context to Docker daemon 9.216kB
Step 1/2 : FROM httpd
latest: Pulling from library/httpd
af107e978371: Pull complete
eba4da411ea0: Pull complete
4f4fb700ef54: Pull complete
ed4d6291a6c2: Pull complete
b42c390e1de9: Pull complete
eafe388a0bb8: Pull complete
Digest: sha256:f0a93744d8006e6f7ee5086c9ddccdcfa33d1091f15269a00547b4c382459c1f
Status: Downloaded newer image for httpd:latest
--> 6fd77d7e5eb7
Step 2/2 : RUN echo "<h1>Hi Team welcome to deployment1-app1</h1>" > /usr/local/apache2/htdocs/index.html
--> Running in 2ec8c7983979
Removing intermediate container 2ec8c7983979
--> 2b4b0bd79968
Successfully built 2b4b0bd79968
```

```
[root@workstation myobjects]# docker login
Login with your Docker ID to push and pull images from Docker Hub. If you don't have a Docker ID, head over to https://hub.docker.com to create one.
Username: shan5a6
Password:
WARNING! Your password will be stored unencrypted in /root/.docker/config.json.
Configure a credential helper to remove this warning. See
https://docs.docker.com/engine/reference/commandline/login/#credentials-store

Login Succeeded
[root@workstation myobjects]# docker push shan5a6/deployment1-app1:v1.0
The push refers to repository [docker.io/shan5a6/deployment1-app1]
9af29645bb07: Pushed
8cd4026118f7: Mounted from library/httpd
8c10599774e2: Mounted from library/httpd
87939ee964f4: Mounted from library/httpd
5f70bf18a086: Mounted from library/httpd
1b65777f1238: Mounted from library/httpd
7292cf786aa8: Mounted from library/httpd
v1.0: digest: sha256:4340efe6d3c2f0befbc58a42bdca7df92275eb485e6dc2623b057fa116c1ad40 size: 1779
[root@workstation myobjects]# vi deploy.yaml
[root@workstation myobjects]# vi Dockerfile
```

Do the same for App2:

App2:

Dockerfile:

```
FROM httpd
RUN echo "<h1>Hi Team welcome to deployment2-app2</h1>" >
/usr/local/apache2/htdocs/index.html
```

```
docker build -t "shan5a6/deployment2-app2:v1.0" .
```

```
[root@workstation myobjects]# vi Dockerfile
[root@workstation myobjects]# cat Dockerfile
FROM httpd
RUN echo "<h1>Hi Team welcome to deployment2-app2</h1>" > /usr/local/apache2/htdocs/index.html
[root@workstation myobjects]# docker build -t "shan5a6/deployment2-app2:v1.0" .
Sending build context to Docker daemon 9.216KB
Step 1/2 : FROM httpd
--> 6fd77d7e5eb7
Step 2/2 : RUN echo "<h1>Hi Team welcome to deployment2-app2</h1>" > /usr/local/apache2/htdocs/index.html
--> Running in 9073cle389fb
Removing intermediate container 9073cle389fb
--> fcdbd7ca8d61c
Successfully built fcdbd7ca8d61c
Successfully tagged shan5a6/deployment2-app2:v1.0
[root@workstation myobjects]# docker push shan5a6/deployment2-app2:v1.0
The push refers to repository [docker.io/shan5a6/deployment2-app2]
2c8cf7350058: Pushed
8cd4026118f7: Mounted from shan5a6/deployment1-app1
8c10599774e2: Mounted from shan5a6/deployment1-app1
87939ee964f4: Mounted from shan5a6/deployment1-app1
5f70bf18a086: Mounted from shan5a6/deployment1-app1
1b65777f1238: Mounted from shan5a6/deployment1-app1
7292cf786aa8: Mounted from shan5a6/deployment1-app1
v1.0: digest: sha256:ad05f72268f31f4ee78c396c4e5106d49f592a98d706dee7865fa168a655aa93 size: 1779
[root@workstation myobjects]# kubectl create deploy deployment1-app1 --image shan5a6/deployment1-app1
```

Create the deployments along with services in javaapp namespace:

```
kubectl create ns javaapp
```

App2:

Deployment:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: app2-deploy
    name: app2-deploy
    namespace: javaapp
spec:
  replicas: 1
  selector:
```

```
matchLabels:  
  app: app2-deploy  
strategy: {}  
template:  
  metadata:  
    creationTimestamp: null  
    labels:  
      app: app2-deploy  
spec:  
  containers:  
    - image: shan5a6/deployment2-app2:v1.0  
      name: app2  
      resources: {}  
status: {}
```

Service:

```
apiVersion: v1  
kind: Service  
metadata:  
  creationTimestamp: null  
  labels:  
    app: app2-deploy  
  name: app2-deploy  
  namespace: javaapp  
spec:  
  ports:  
    - port: 80  
      protocol: TCP  
      targetPort: 80  
  selector:  
    app: app2-deploy  
status:  
  loadBalancer: {}
```

Same goes with APP1:

Deployment:

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  creationTimestamp: null  
  labels:  
    app: app1-deploy
```

```
name: app1-deploy
namespace: javaapp
spec:
replicas: 1
selector:
  matchLabels:
    app: app1-deploy
strategy: {}
template:
  metadata:
    creationTimestamp: null
    labels:
      app: app1-deploy
  spec:
    containers:
      - image: shan5a6/deployment1-app1:v1.01
        name: app1
        resources: {}
status: {}
```

Service:

```
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    app: app1-deploy
  name: app1-deploy
  namespace: javaapp
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
  selector:
    app: app1-deploy
status:
  loadBalancer: {}
```

Let's configure the ingress rules:

```
[root@workstation ~]# cat app-ingress.yaml
```

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: application-access
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
    nginx.ingress.kubernetes.io/use-regex: "true"
spec:
  ingressClassName: nginx
  rules:
  - http:
      paths:
      - path: /app1
        pathType: Prefix
        backend:
          service:
            name: app1-deploy
            port:
              number: 80
      - path: /app2
        pathType: Prefix
        backend:
          service:
            name: app2-deploy
            port:
              number: 80
      - path: /
        pathType: Prefix
        backend:
          service:
            name: app1-deploy
            port:
              number: 80
```

```
[root@workstation myobjects]# cat ingress.yaml
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: application-access
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
    nginx.ingress.kubernetes.io/use-regex: "true"
spec:
  ingressClassName: nginx
  rules:
  - http:
      paths:
      - path: /appl
        pathType: Prefix
        backend:
          service:
            name: deployment1-app1-svc1
            port:
              number: 80
      - path: /app2
        pathType: Prefix
        backend:
          service:
            name: deployment2-app2-svc2
            port:
              number: 80
      - path: /
        pathType: Prefix
        backend:
```

```
[root@workstation myobjects]# kubectl delete -A ValidatingWebhookConfiguration ingress-nginx-admission
validatingwebhookconfiguration.admissionregistration.k8s.io "ingress-nginx-admission" deleted
[root@workstation myobjects]#
[root@workstation myobjects]# kubectl apply -f ingress.yaml -n javaapp
ingress.networking.k8s.io/application-access created
[root@workstation myobjects]#
[root@workstation myobjects]# kubectl get ingress -n javaapp
NAME                CLASS      HOSTS   ADDRESS
application-access  nginx     *       ad7fbf9075cf4c9e90a252cda89978e-54684905.us-east-2.elb.amazonaws.com
[root@workstation myobjects]# kubectl get svc -n javaapp
NAME         TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
kubernetes   ClusterIP   10.96.0.1    <none>        443/TCP   19s
```

Execute the below if you are facing any issues during ingress definition creation.

kubectl delete -A ValidatingWebhookConfiguration ingress-nginx-admission

Final verification:

Not secure ad7fbf9075cf4c9e90a252cda89978e-54684905.us-east-2.elb.amazonaws.com

Hi Team welcome to deployment1-app1

← → ⌂ Not secure ad7fbf9075cf4c9e90a252cda89978e-54684905.us-east-2.elb.amazonaws.com/app1

Hi Team welcome to deployment1-app1

← → ⌂ Not secure ad7fbf9075cf4c9e90a252cda89978e-54684905.us-east-2.elb.amazonaws.com/app2

Hi Team welcome to deployment2-app2

DVS Technologies